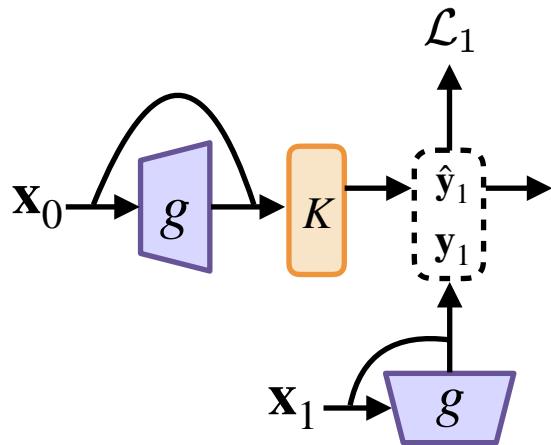


Project Summary W1-2

Recap: Two Different NNs

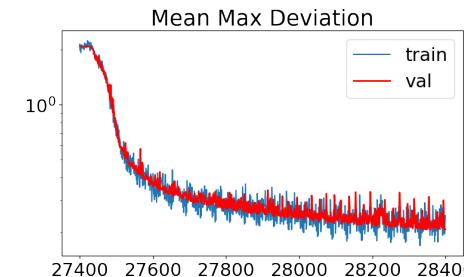
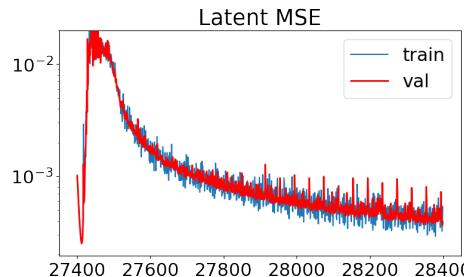
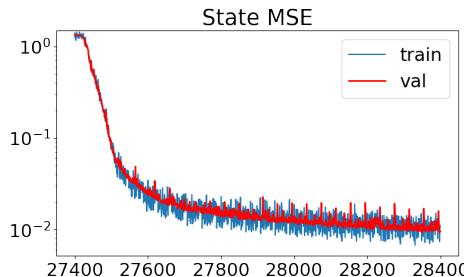
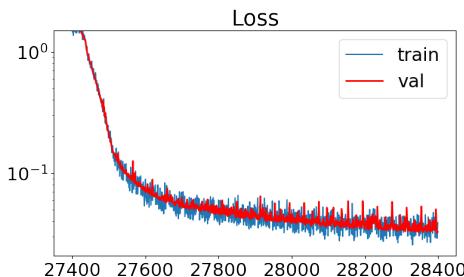
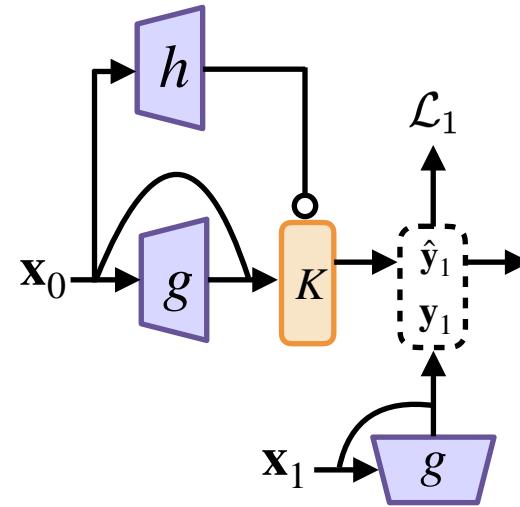
LREN:

Linearly Recurrent Encoder Network



DENIS:

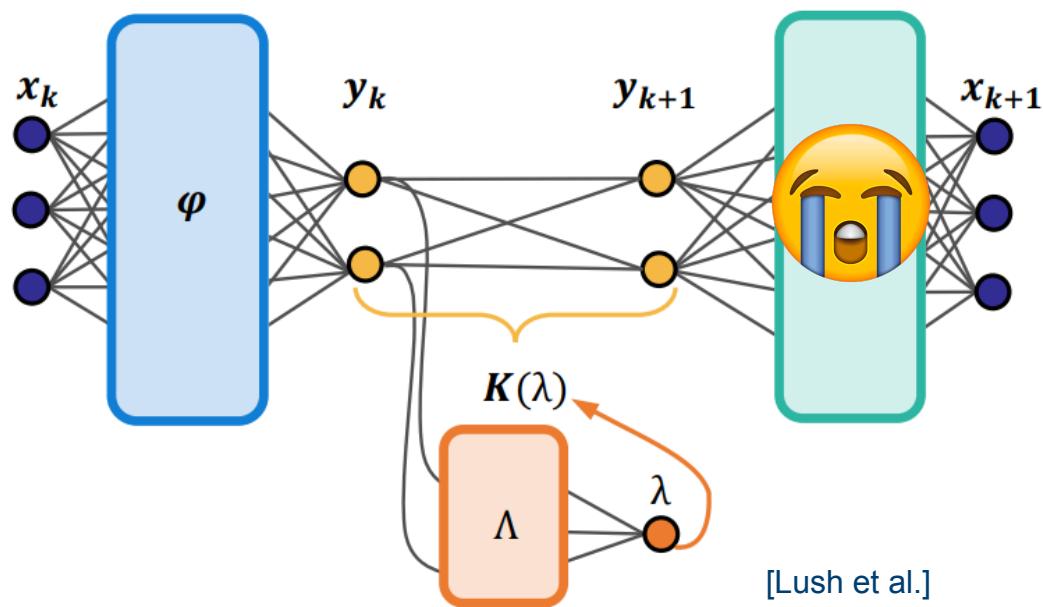
Deep Encoder Network with Initial State Parameterisation



Side Note

Ideally, only 2 latent layers

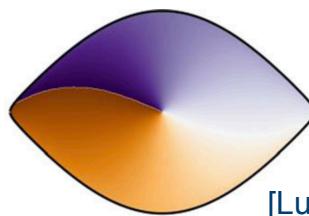
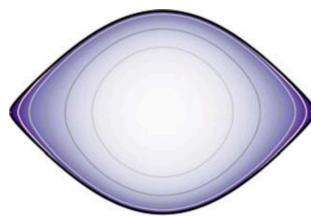
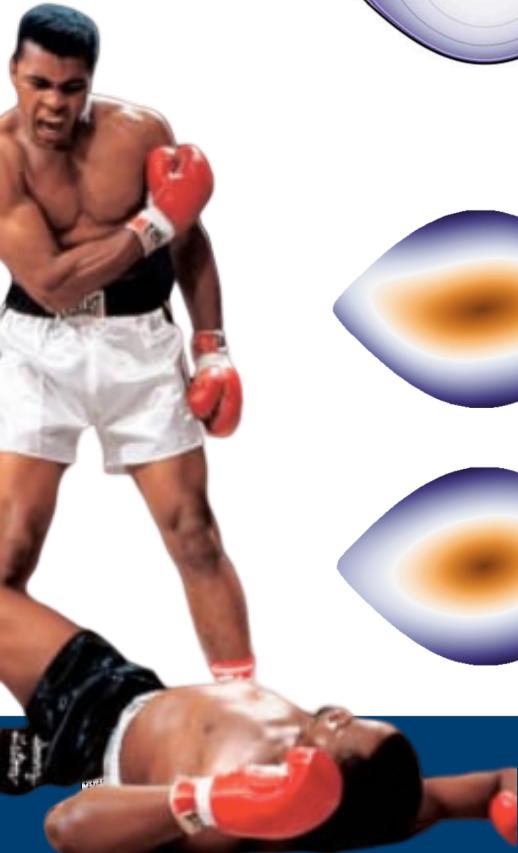
$$\text{Dim}(y) >> 2$$



[Lush et al.]

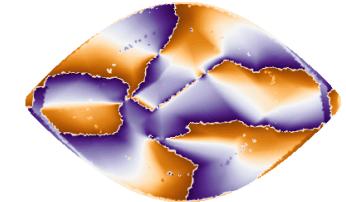
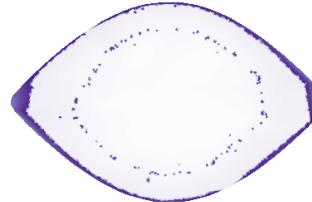
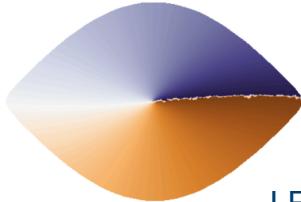
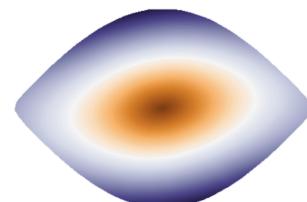
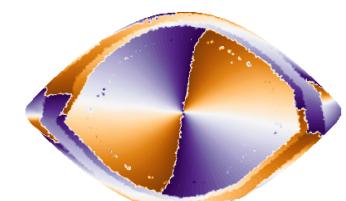
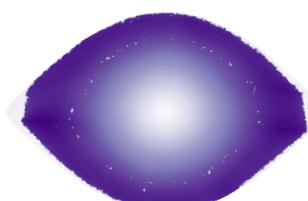
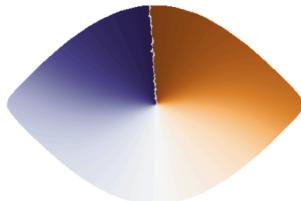
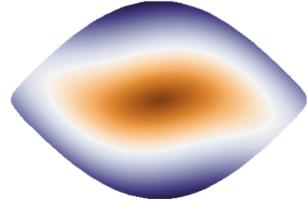
DENIS VS LREN Showdown

Comparison of eigenfunctions:



$$E = m \frac{(l\dot{\theta})^2}{2} + mgl(1 - \cos(\theta))$$

[Lush et al.]

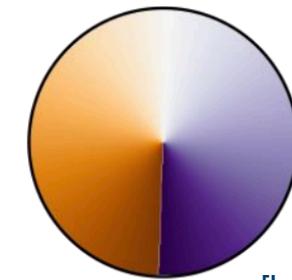
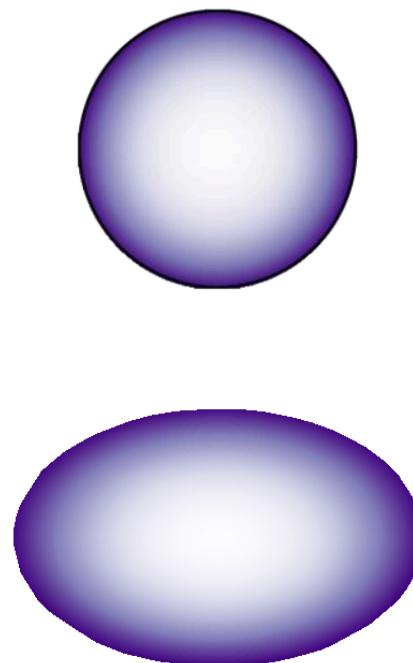
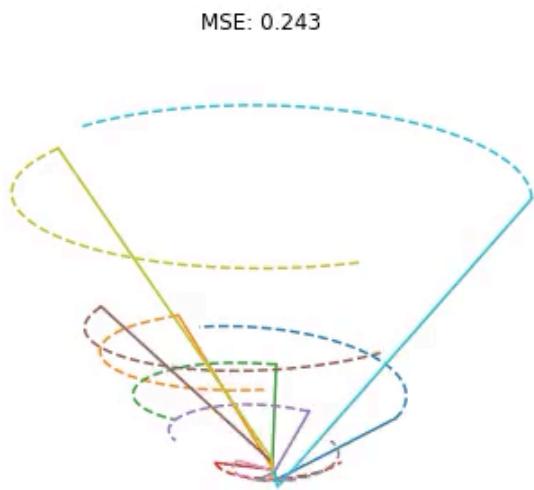


LREN

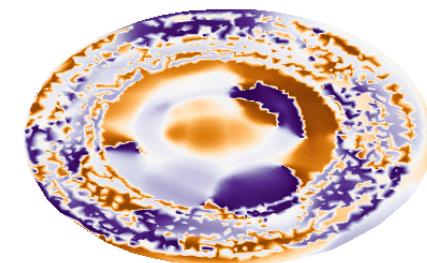
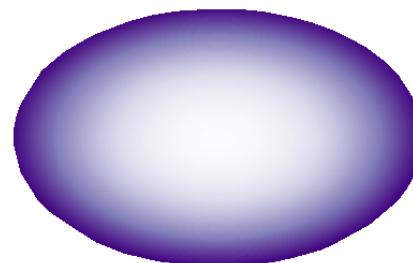
DENIS

Side Note: DENIS Fluid Example

Comparison of eigenfunctions:



[Lush et al.]

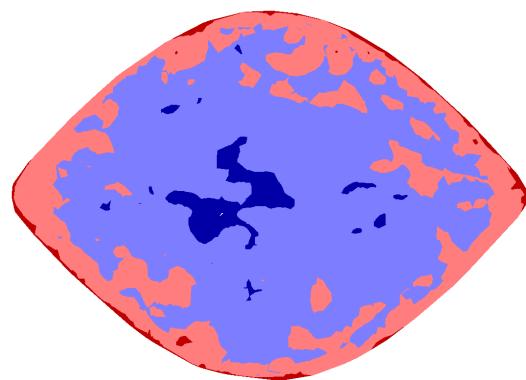


DENIS

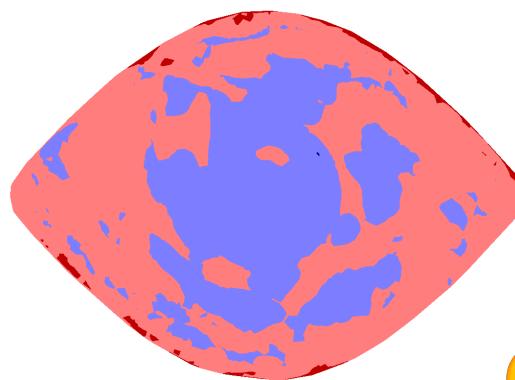
DENIS VS LREN Showdown

Do we get better predictions?

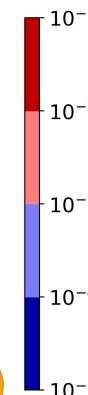
DENIS



LREN



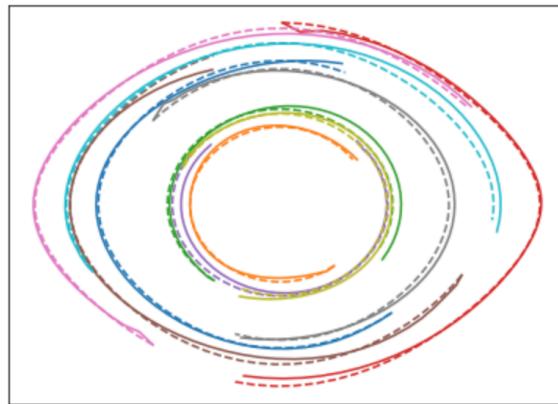
MSE



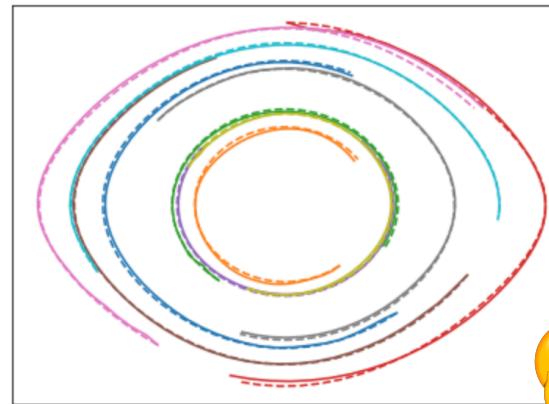
Average prediction
error across all 5
seconds



LREN



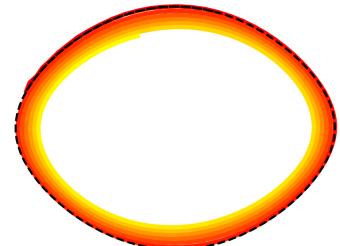
DENIS



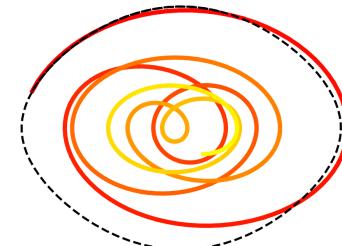
DENIS VS LREN Showdown

Random state initialisations

DENIS



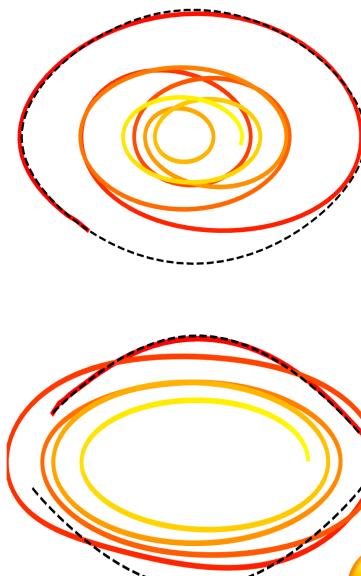
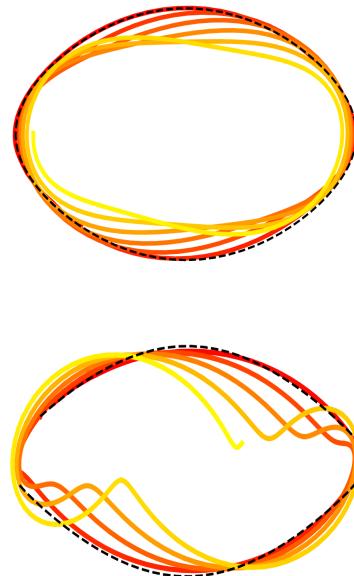
LREN



How about long time-horizon predictions?
Setting prediction horizon to 50 seconds (500 steps!)

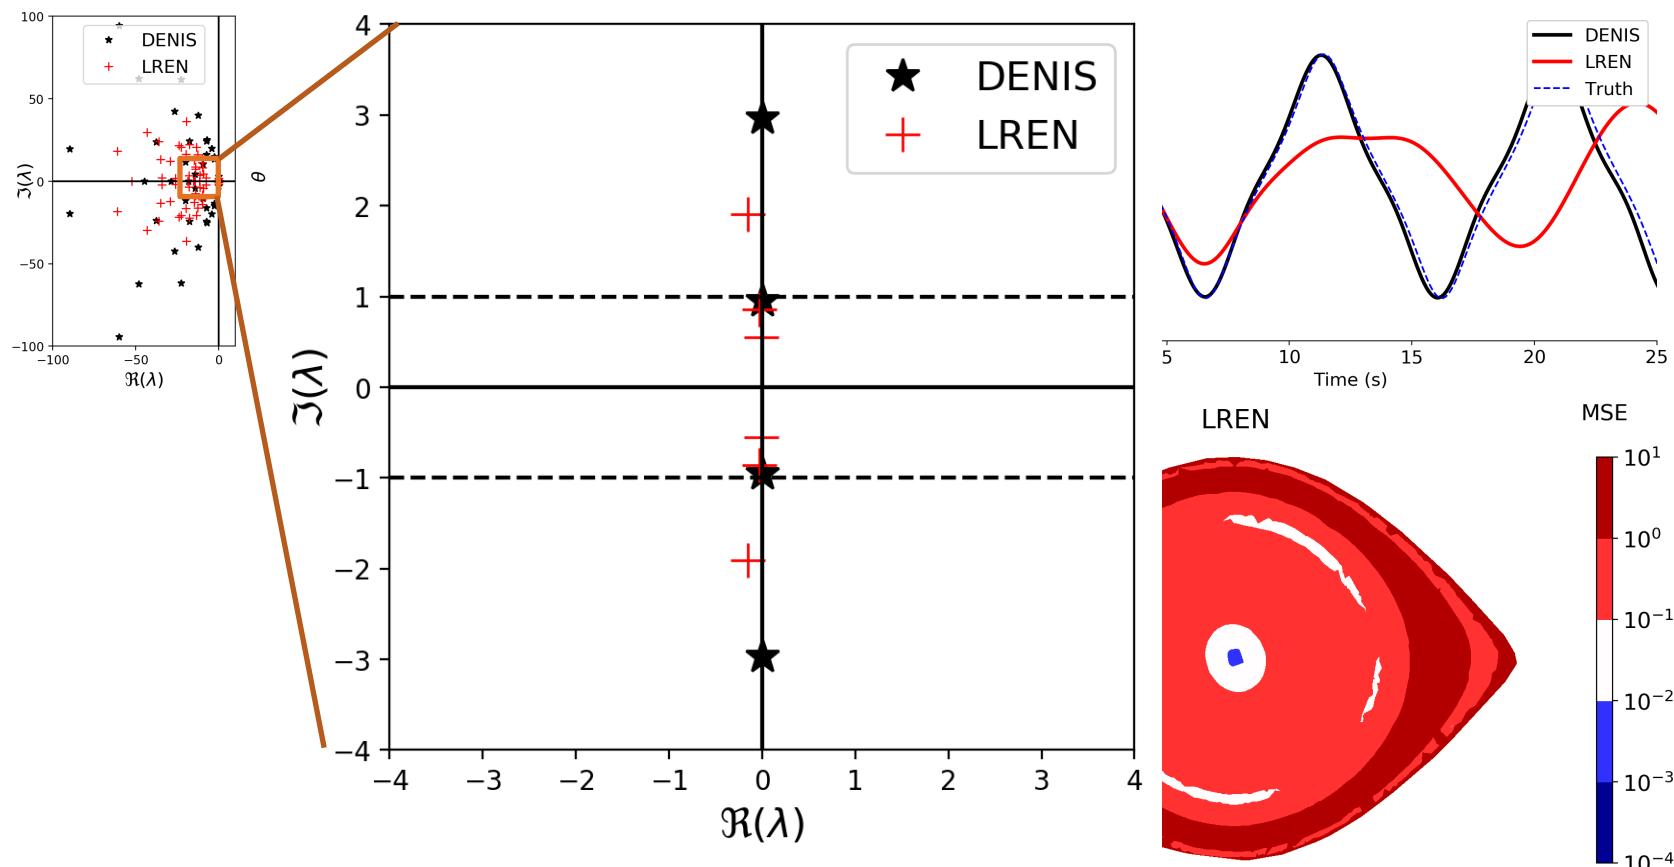
t=50

t=0

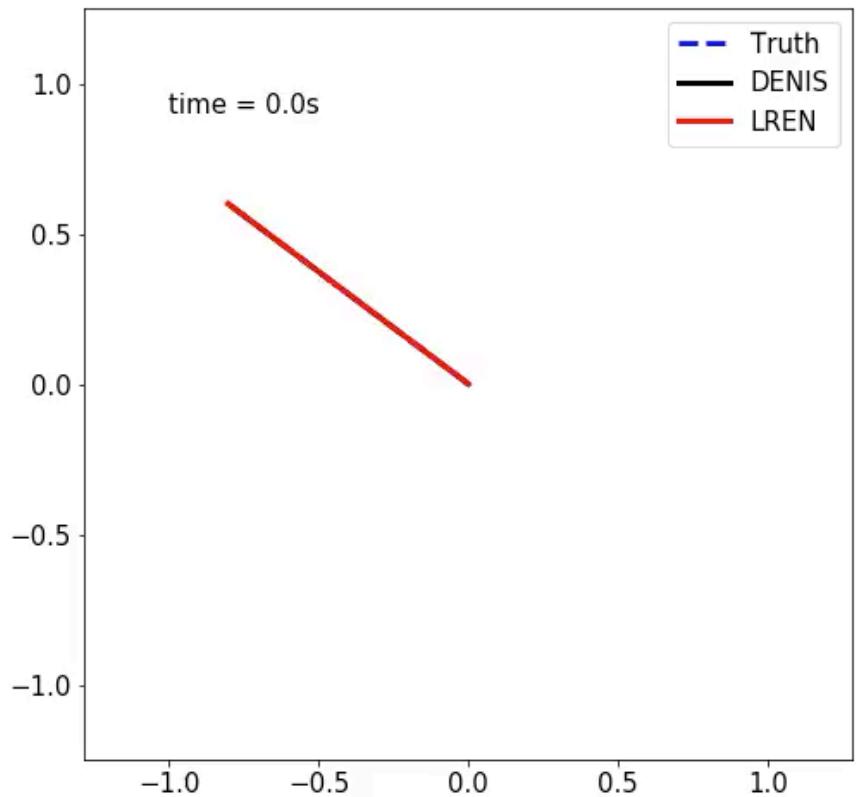
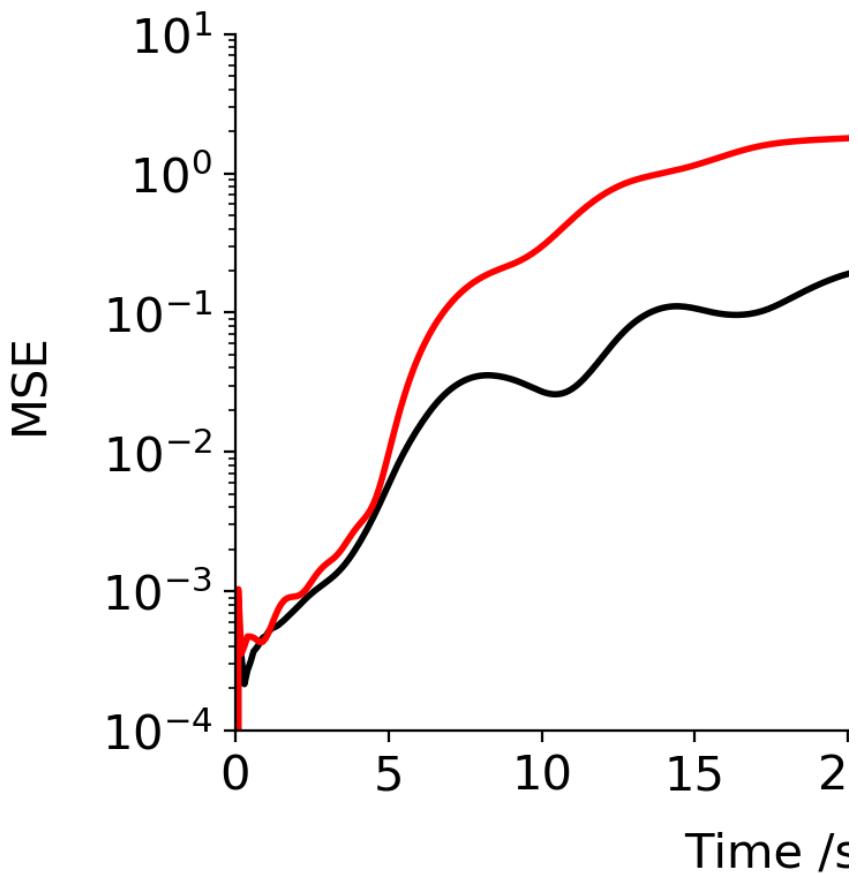


DENIS VS LREN Showdown

How about long time-horizon predictions?
Setting prediction horizon to 50 seconds (500 steps!)

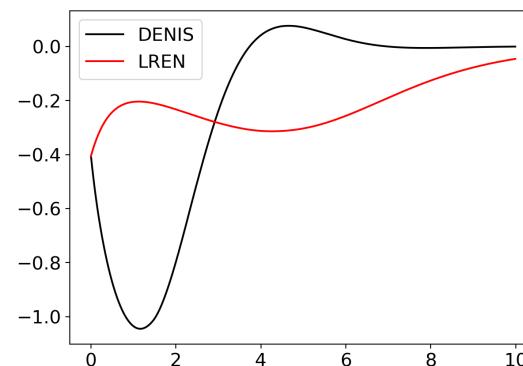
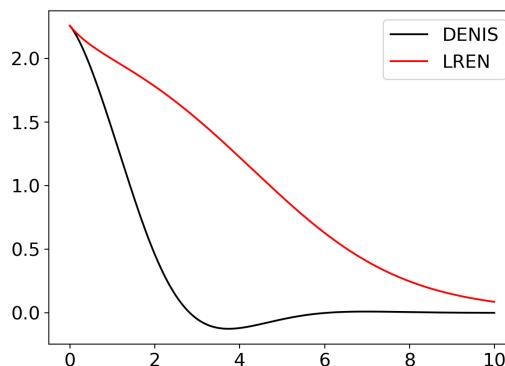
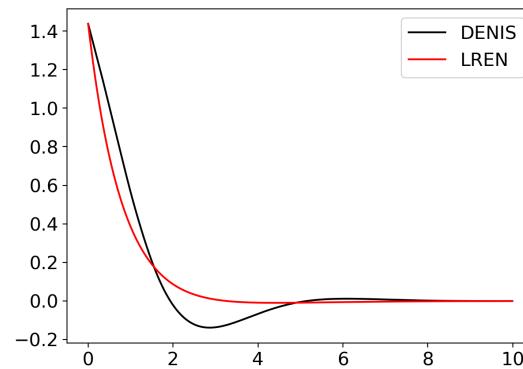
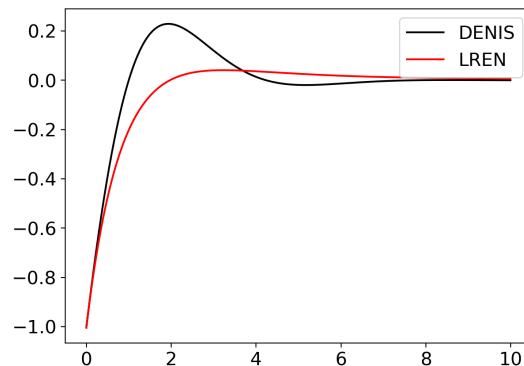


DENIS VS LREN Showdown



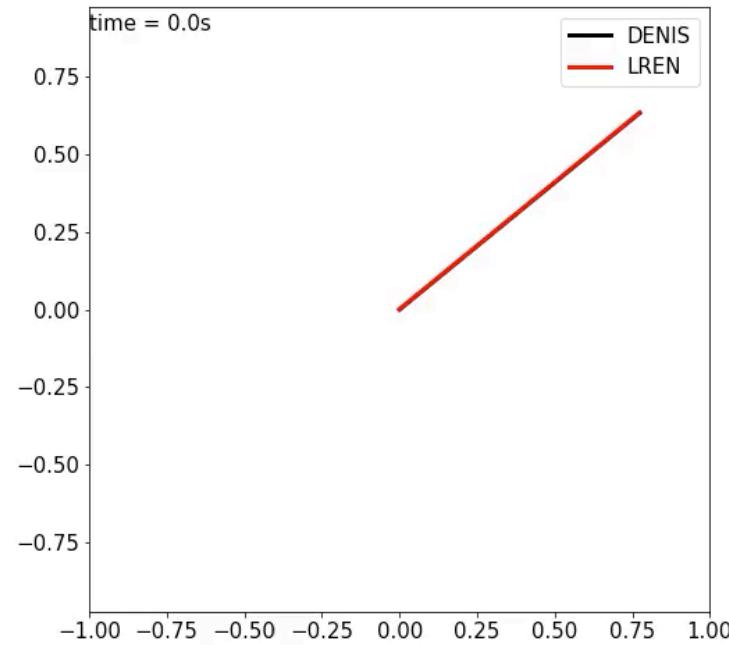
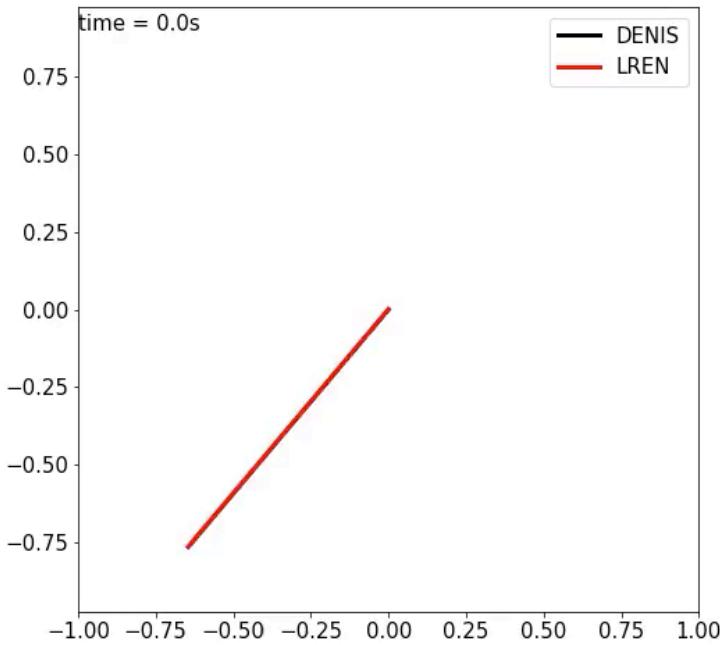
DENIS VS LREN Showdown

How about KOOC? - Same as before, LQR on the latent dims, actual system dynamics



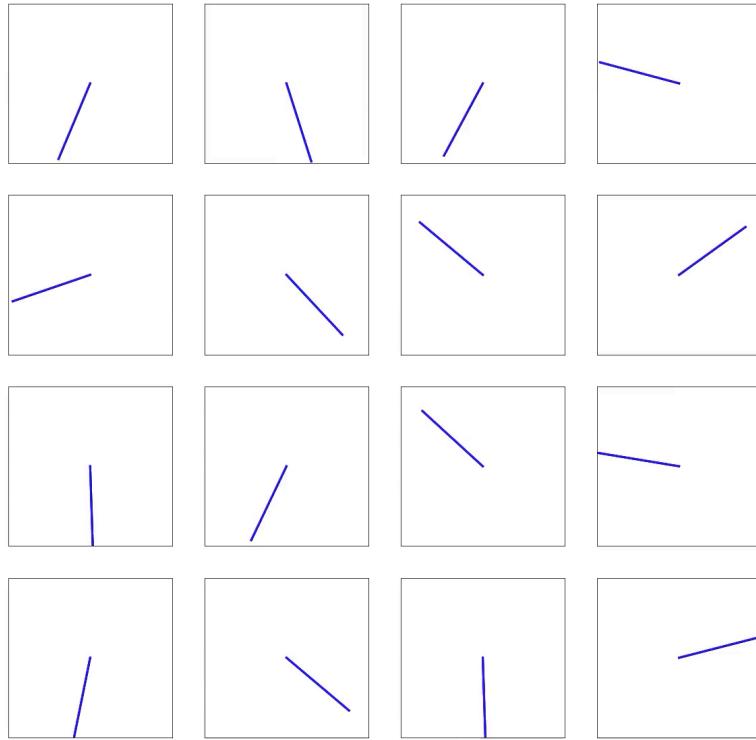
DENIS VS LREN Showdown

How about KOOC? - Same as before, LQR on the latent dims, actual system dynamics



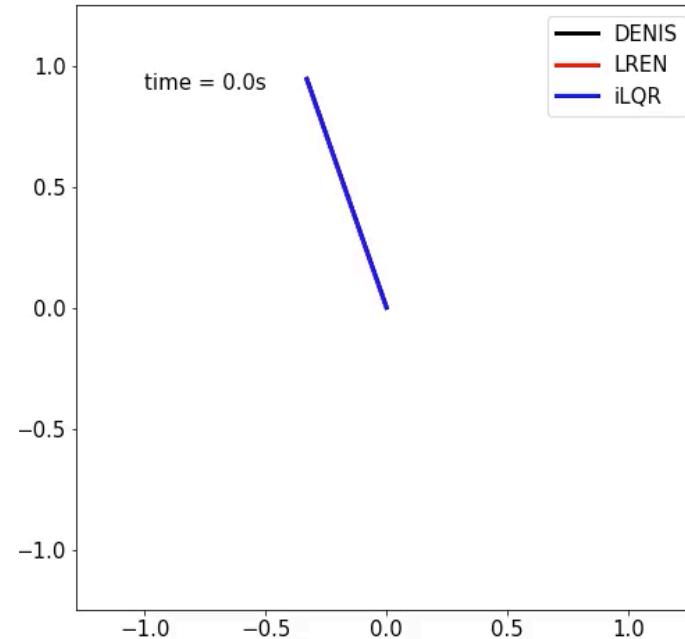
DENIS VS LREN Showdown

— DENIS
— LREN
— iLQR



erative LQR (iLQR)?

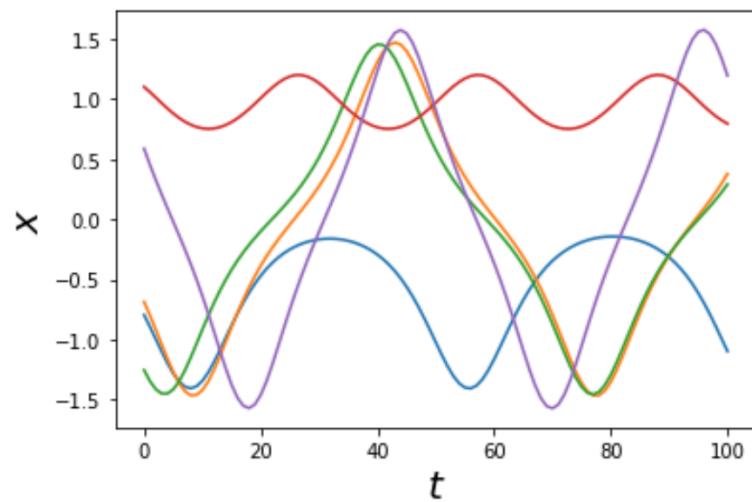
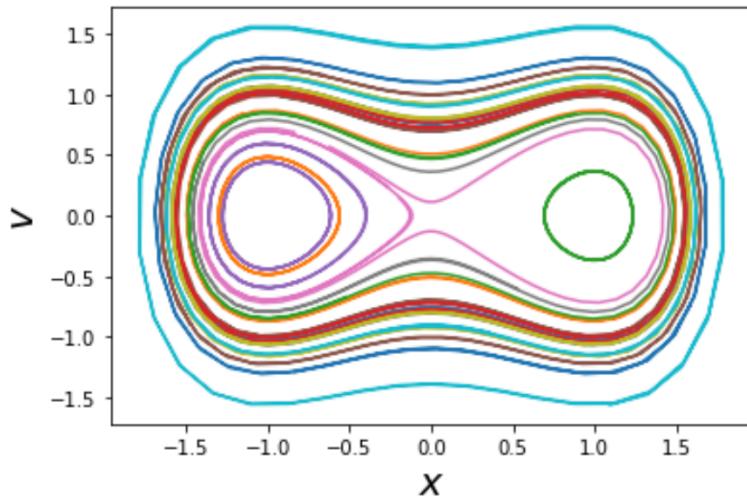
Example: LREN blows up under control input



Next Step: Duffing Oscillator & New Architecture?

$$\dot{\mathbf{x}} = \begin{bmatrix} x_2 \\ x_1 - x_1^3 \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} u$$

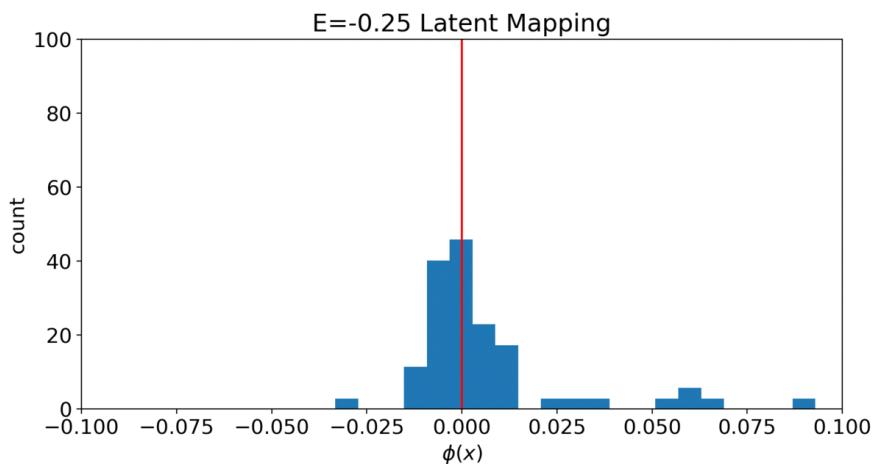
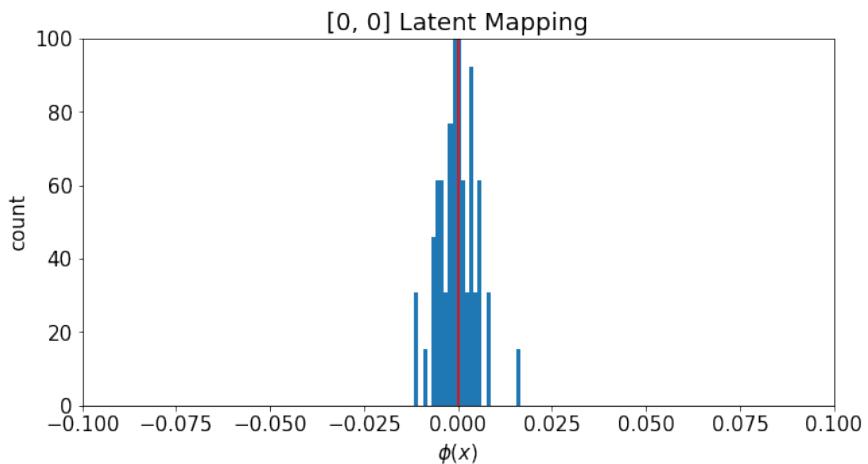
$$\mathcal{H} = \frac{1}{2}x_2^2 - \frac{1}{2}x_1^2 + \frac{1}{4}x_1^4$$



Project Summary W3-4

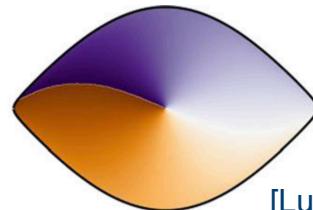
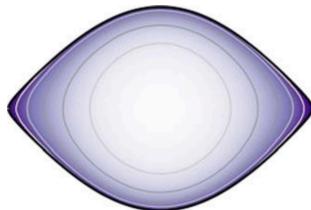
1. Does $\Phi(0)=0$?

- Answer: Not really.



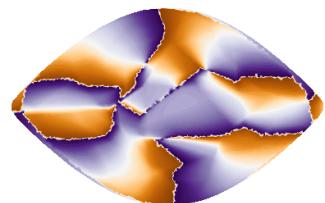
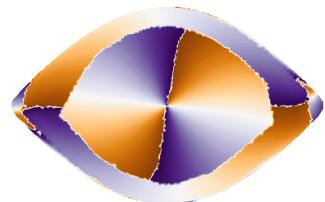
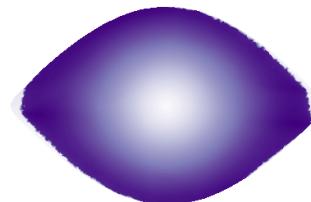
Discontinuity Issues

- Not easy to solve :(

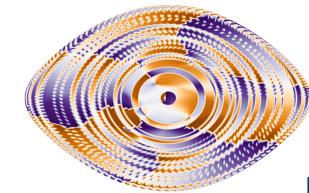
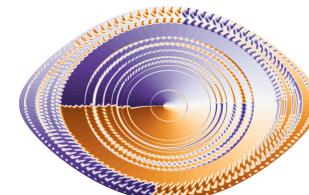
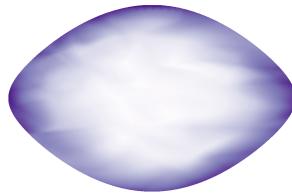
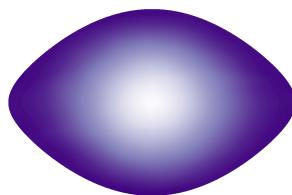


$$E = m \frac{(l\dot{\theta})^2}{2} + mgl(1 - \cos(\theta))$$

[Lush et al.]



DENIS

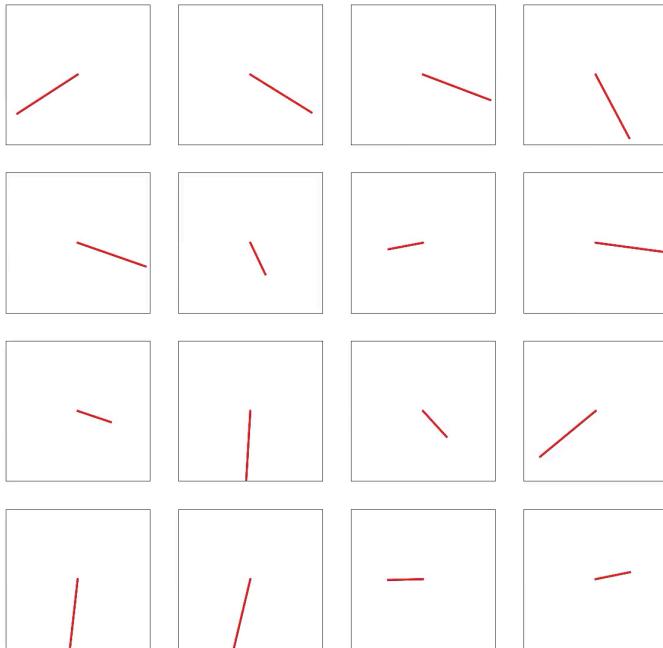


DENIS,
smoothed

MPC?

- Done, but bugs? Parameters? Still not there yet....

— KLQR
— iLQR
— KMPC

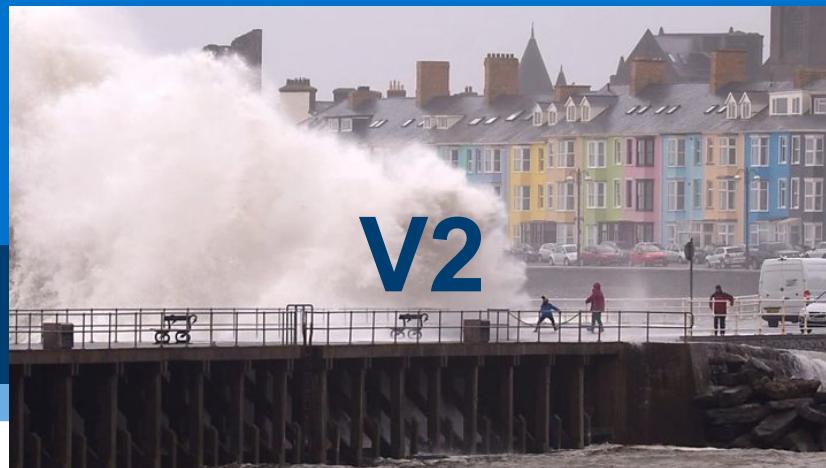


— KLQR
— iLQR
— KMPC



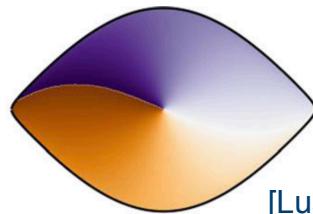
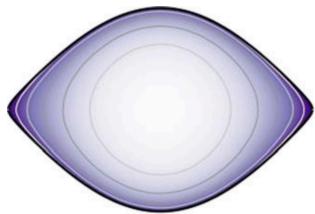
$$\begin{aligned} & \text{minimize}_{U \in \mathbb{R}^{m \times N_p}} U^\top H U + h^\top U + z_0^\top G U \\ & \text{subject to } L U + M z_0 \leq c \\ & \text{parameter } z_0 = \psi(x_k) \end{aligned}$$

Project Summary W4-5

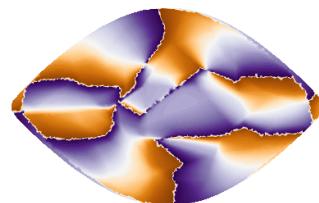
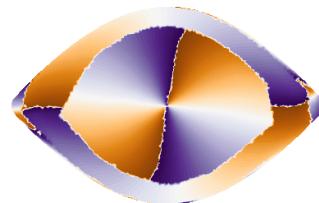
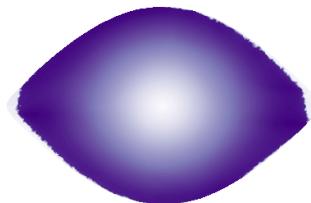
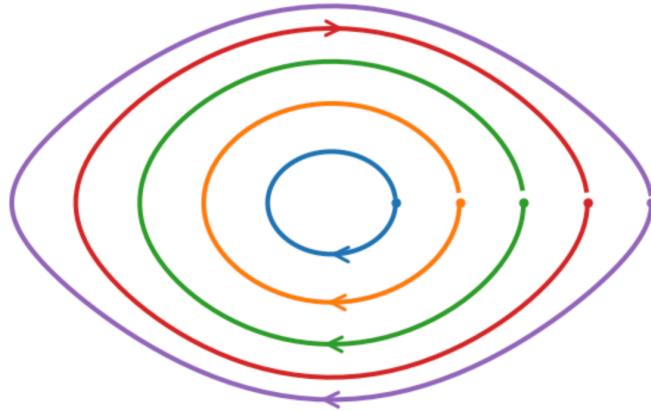


1. Discontinuity Issues

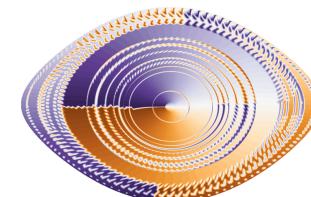
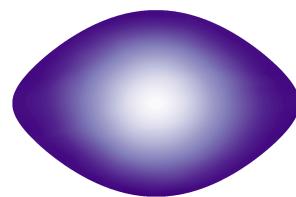
- Behold...



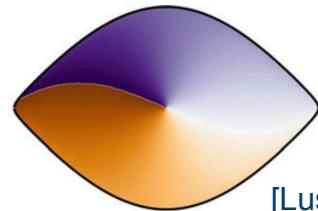
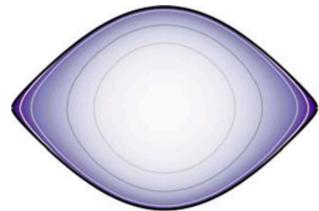
[Lush et al.]



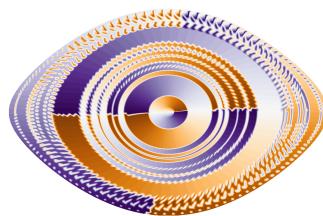
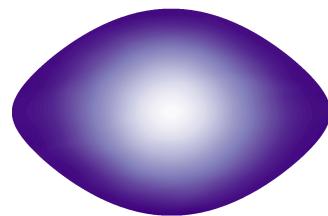
DENIS



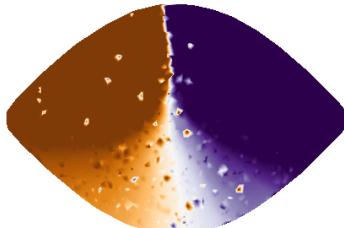
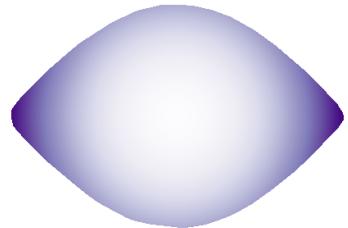
DENIS,
smoothed



[Lush et al.]



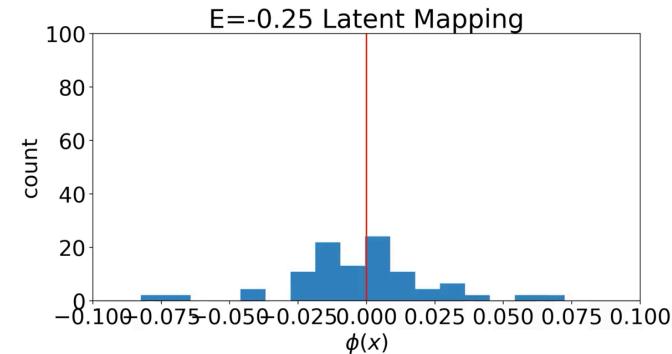
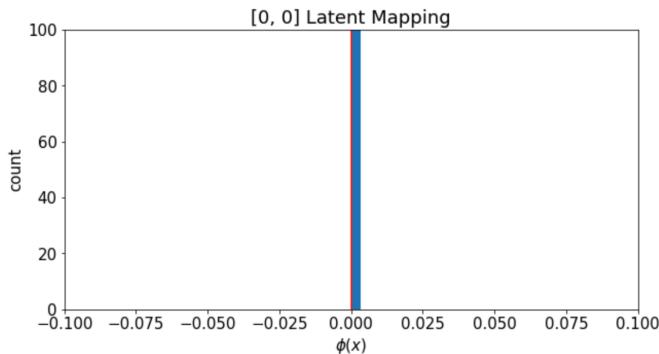
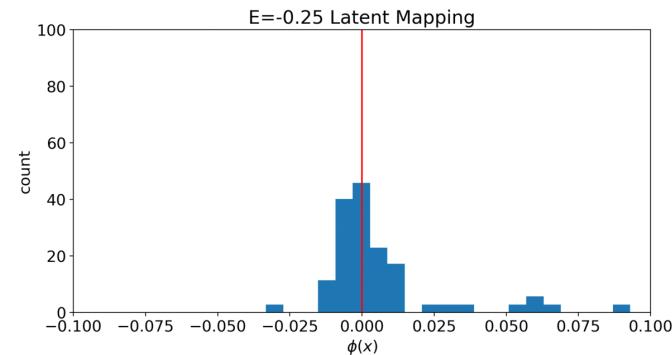
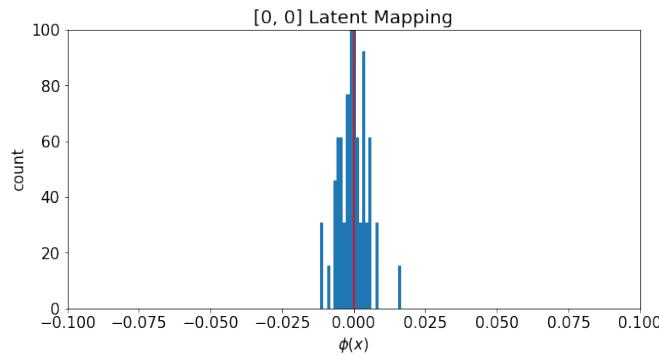
DENIS, previous



DENIS, Now

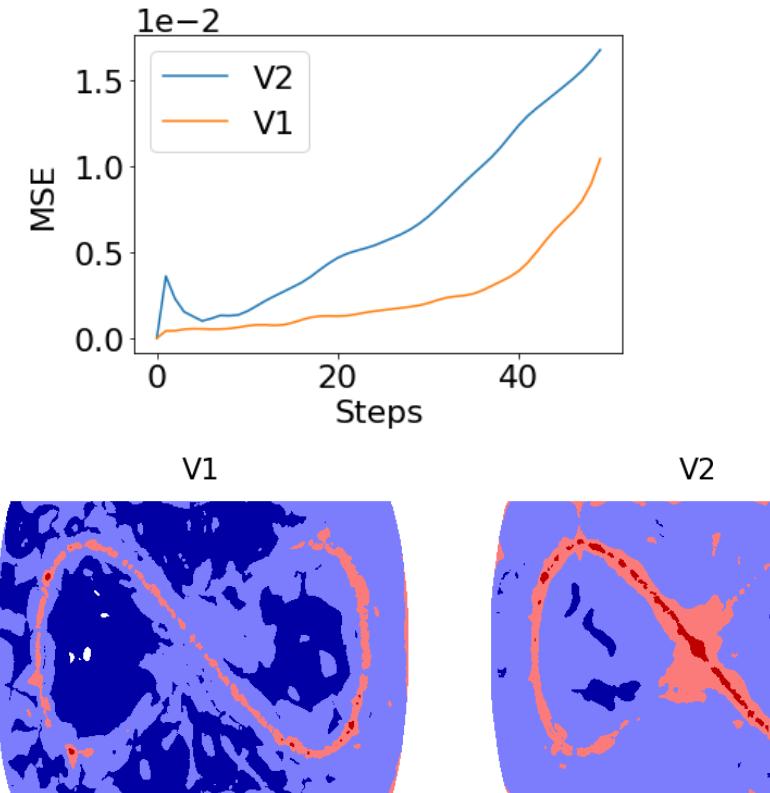
2. Does $\Phi(0)=0$?

- Answer: Yes $\mathcal{L}_{\text{new}} = \mathcal{L}_{\text{old}} + \alpha \frac{1}{N} \sum_n \|K_0^n \phi(\mathbf{0})\|_2$



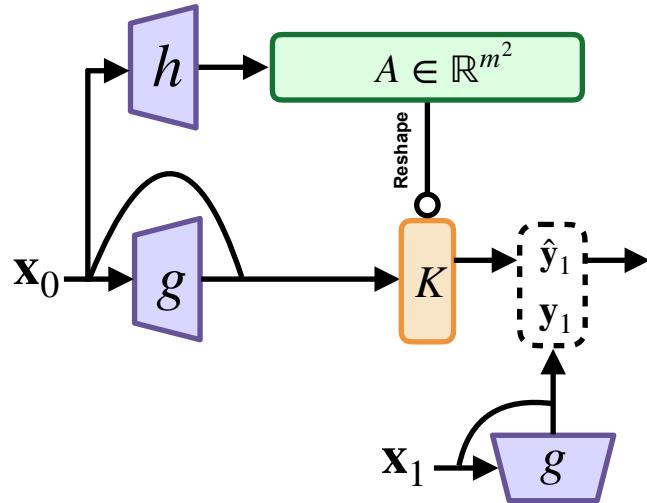
3. Now $\Phi(0)=0$, is MPC/LQR Better?

- Not really...



3. Curse of Dimensionality And Discontinuity

- If $\phi(x) \in \mathbb{R}^m$ then the auxiliary network has $\mathcal{O}(m^2)$ parameters....
- Sat $x \in \mathbb{R}^N$ and let $m = 50N$, network will now have $\mathcal{O}(2500N^2)$ parameters :-o
- Current network has 120K parameters just for a 2D problem
- Discontinuity as both h and g may work together to map eigenfunctions to different locations
- Is there a way to 1. Reduce number of parameters 2. Learn the eigenfunctions



- Jacobian form is not possible - we include our state :(
- Is there a way to globally align the eigenvectors?
- Solutions:

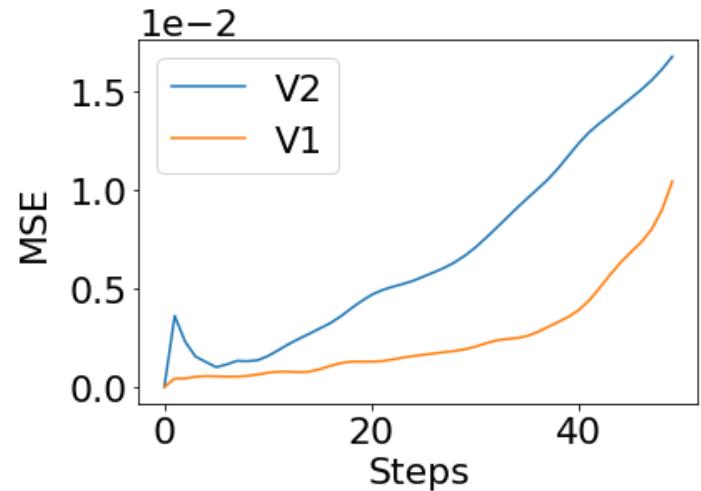
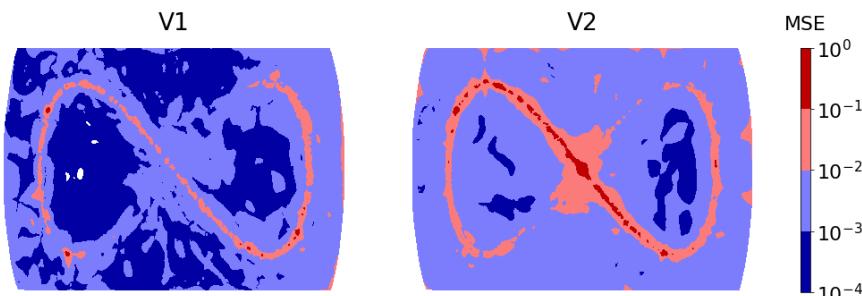
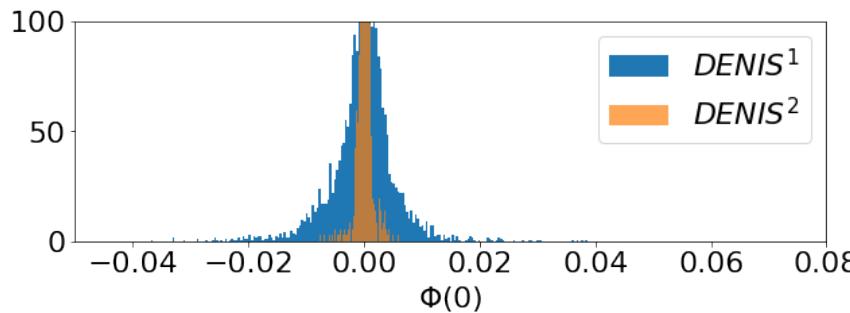
Project Summary W4-5



Wrapping up: Denis's Journey - V2

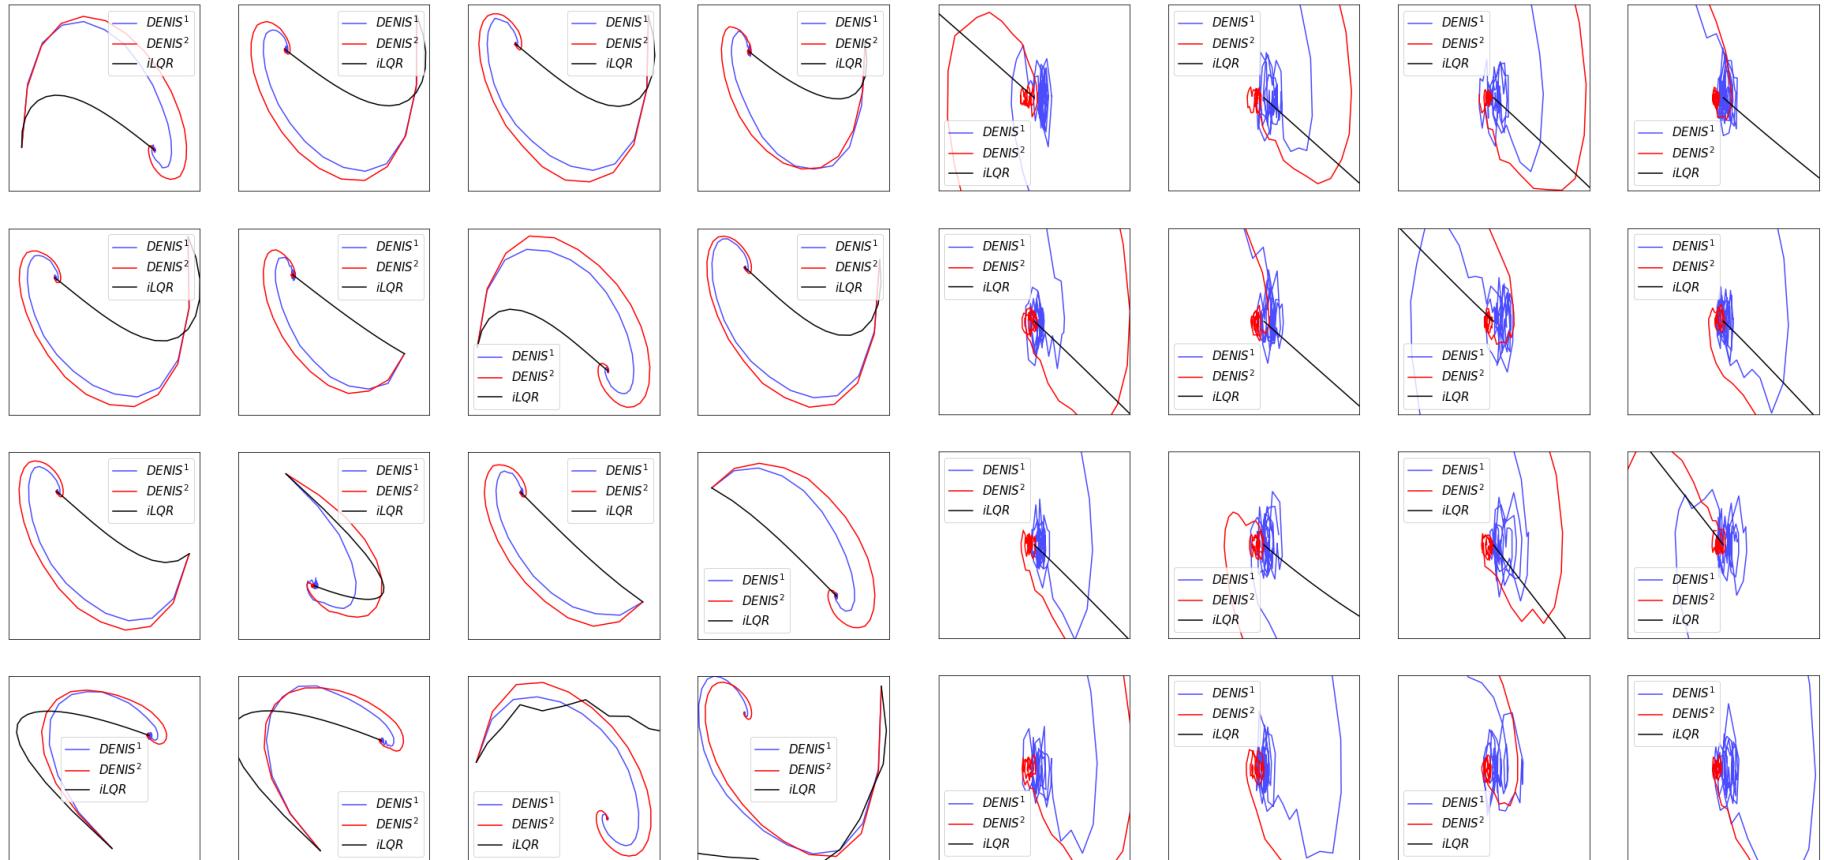
V1 VS V2: From last session: what is the benefit of adding the loss term?

$$\mathcal{L}_{\text{new}} = \mathcal{L}_{\text{old}} + \alpha \frac{1}{N} \sum_n \|K_0^n \phi(\mathbf{0})\|_2$$



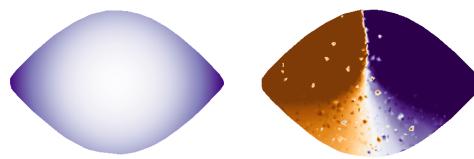
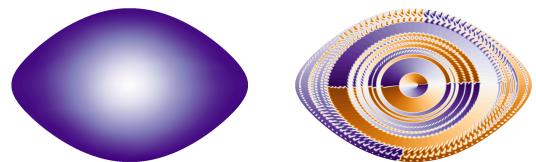
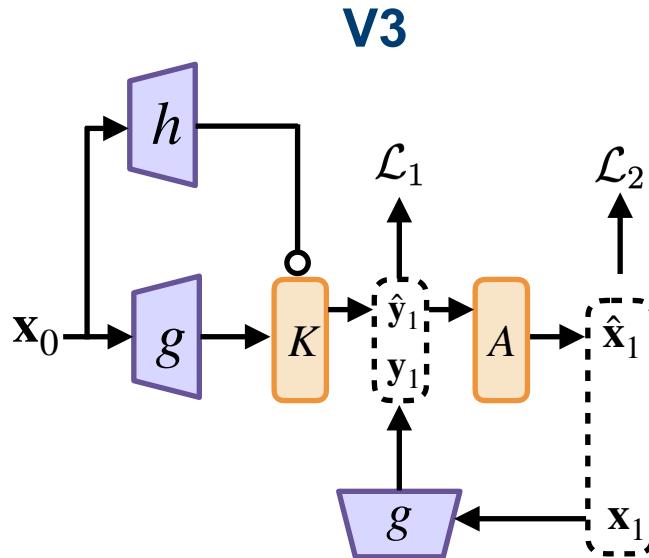
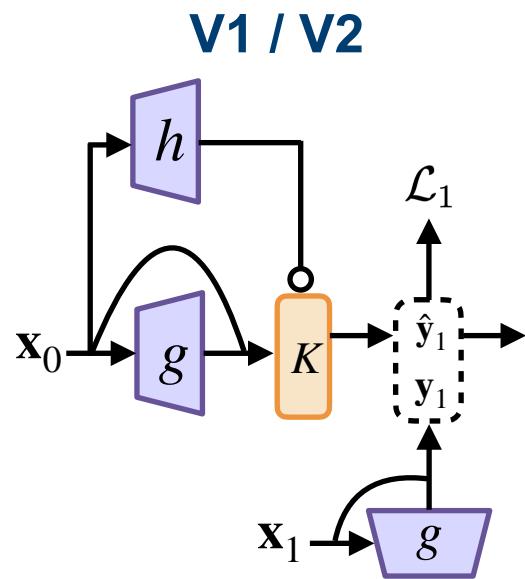
Wrapping up: Denis's Journey

V1 VS V2: From last session: what is the benefit of adding the loss term?

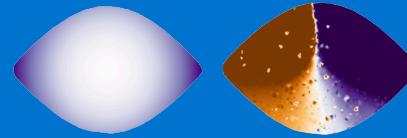


Wrapping up: Denis's Journey -V3

Motivation: Failed to smooth plots, is there a way to ‘guide/force’ the network towards learning the energy function while keeping continuity?



Wrapping up: Denis's Journey

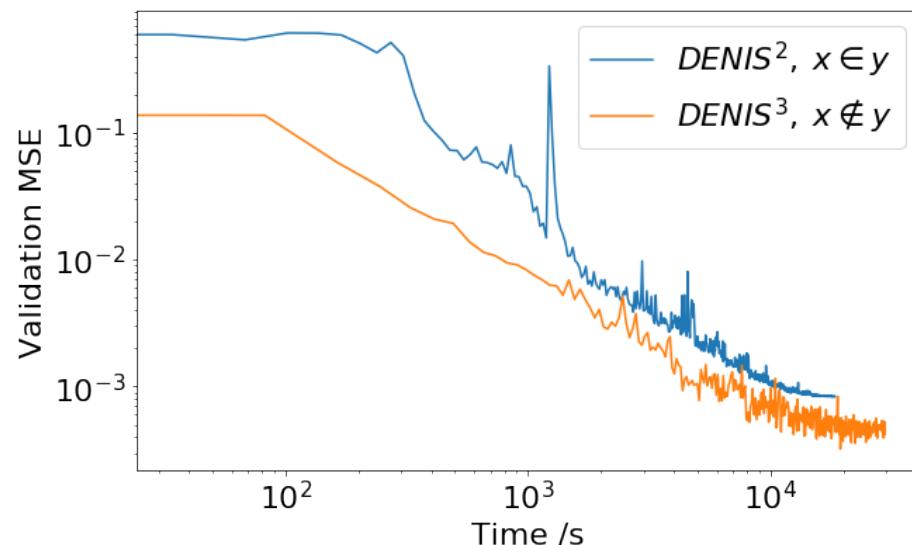


V3: Last attempt on discovering the energy function

$$K = \begin{bmatrix} B_0 & & & \\ & B_1 & & \\ & & B_2 & \\ & \ddots & & \ddots \\ & & & B_N \end{bmatrix}$$
$$B_N = \exp(\mu_n \delta t) \begin{bmatrix} \cos(\omega_n \delta t) & -\sin(\omega_n \delta t) \\ \sin(\omega_n \delta t) & \cos(\omega_n \delta t) \end{bmatrix}$$
$$\omega_n = 0 \quad \forall n \in \{1, \dots, N\}$$

What is its performance and
How is it to train?

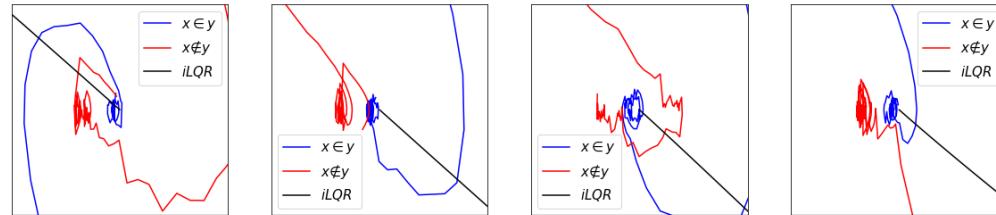
Both approximately same number of
parameters ($\sim 100K$),
Both latent dimension is 150



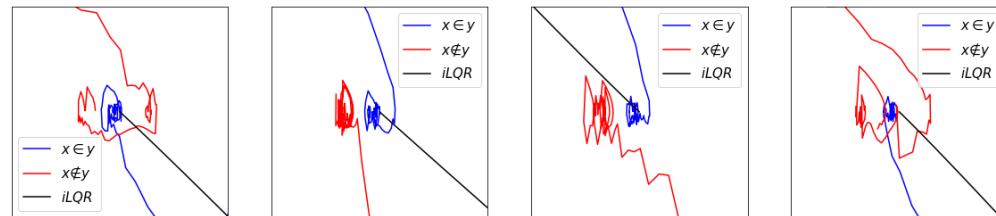
Wrapping up: Denis's Journey

V3: How about MPC?

$$x^+ = f(x) + Bu$$



$$y^+ = Ky + \tilde{B}u$$



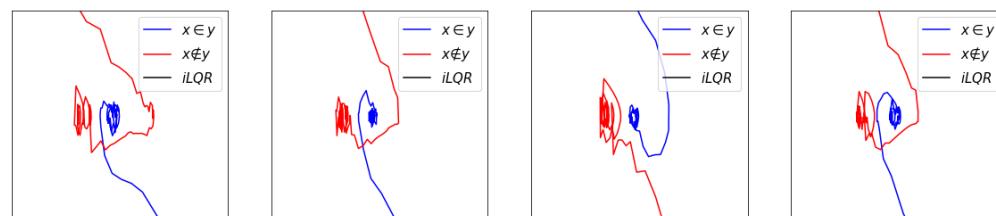
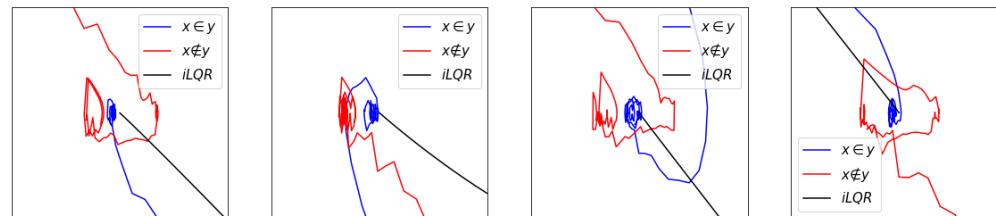
$$y_0 = \Phi(x_0)$$

$$\hat{x} = Ay$$

Easy to show:

$$\tilde{B} = A^\dagger B$$

$$\tilde{Q} = A^T Q A$$



0.1

What have we learnt?

- $\Phi(0) = 0$ is ideal if we want to drive system to 0 and stay there, if this is violated, we may get noisy behaviour at 0. However, we may need to trade off model performance
- $x \in y$ is slightly worse than $x \notin y$ in terms training performance, control wise is uncertain
- Reconstructing the energy function is difficult without a non-linear decoder, but does it really matter?
- The idea behind DENIS is especially useful for MPC
- MPC is pretty good

Introducing: DEINA

Double Encoder for Input Non-Affine Systems

Theory reminder:

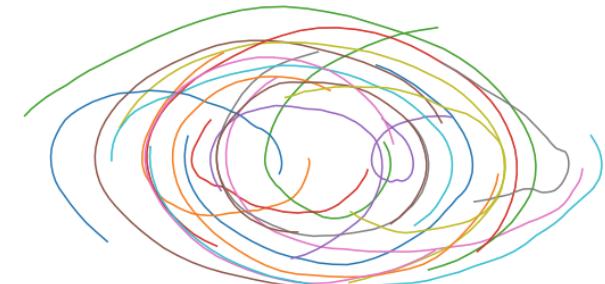
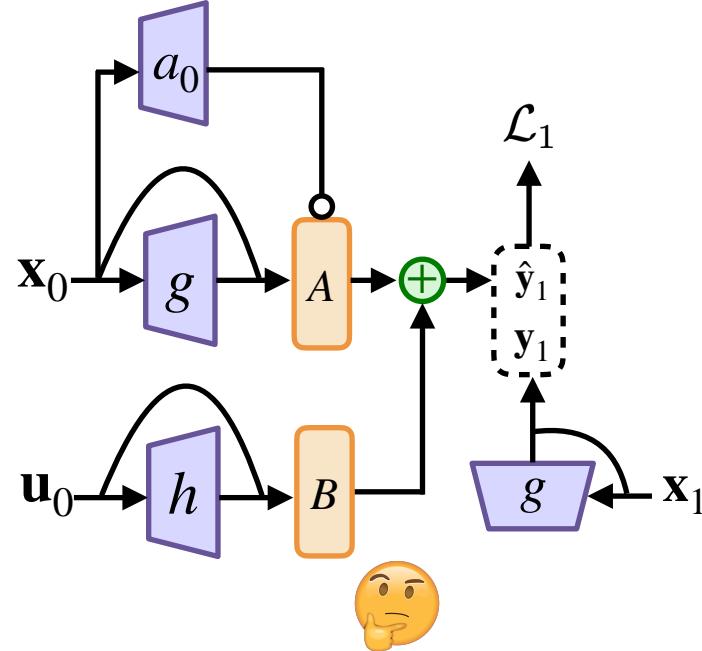
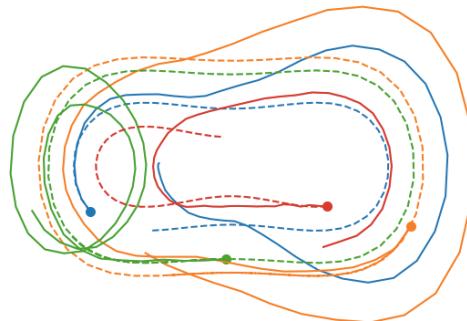
$$\mathbf{x}_{k+1} = f(\mathbf{x}_k, \mathbf{u}_k)$$

$$\mathbf{y}_0 = \Phi_x(\mathbf{x}_0)$$

$$\mathbf{q} = \Phi_u(\mathbf{u})$$

$$\mathbf{y}_{k+1} = A\mathbf{y}_k + B[\mathbf{u} \quad \mathbf{q}]^T$$

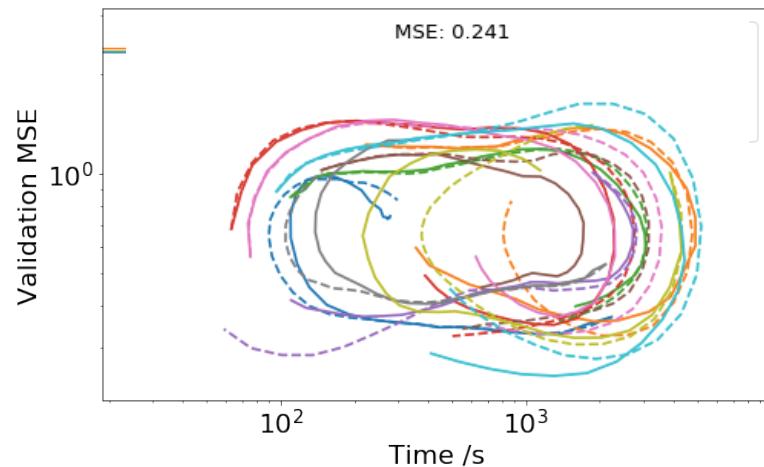
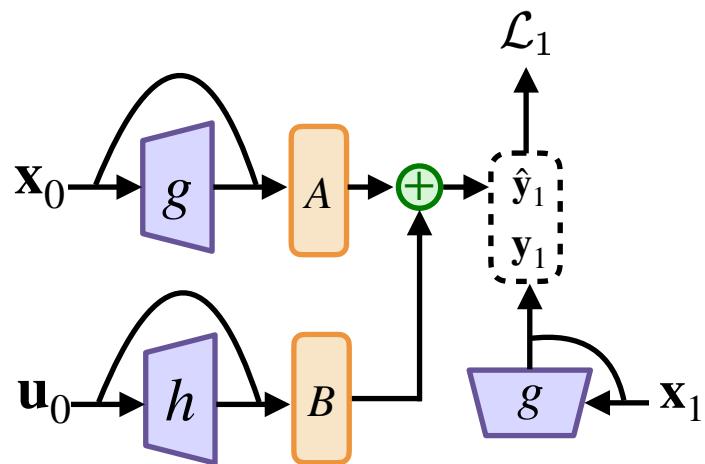
[Mezic]



Initial Results - TL;DR Not Good

Questions:

1. Can we train only with dynamics with input, will the network generalise to 0 input dynamics?
2. Do we parameterise? If so, what information do we give to the auxiliary?
3. What type of input noise do we provide? (Brownian/ Gaussian/ Uniform, etc?)

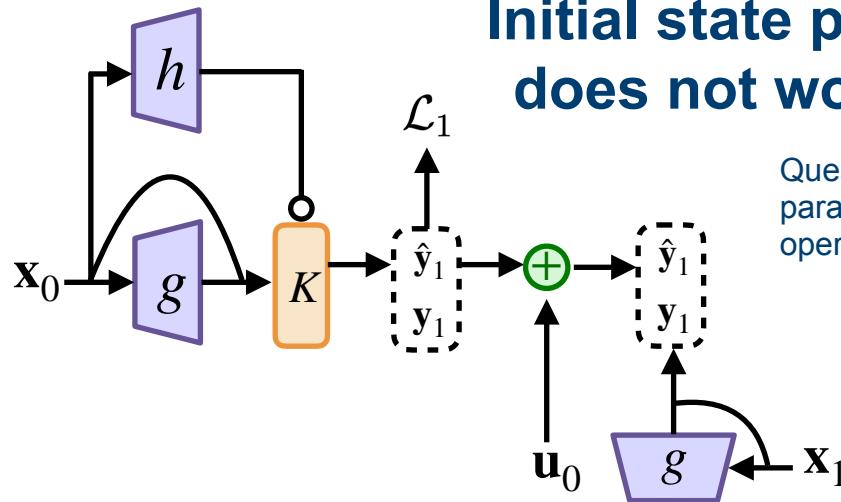


Input Affine: Trivial Extension from DENIS?

Not so much....

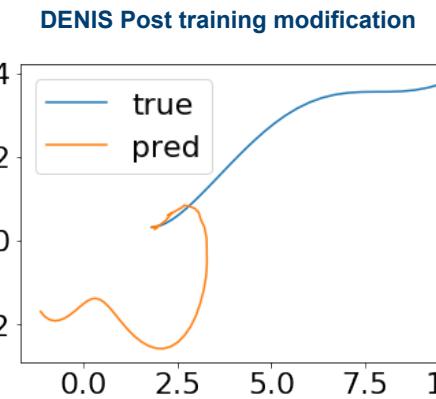
$$\mathbf{y}^1 = A\mathbf{y}^0 + B\mathbf{u}^0$$

$$\mathbf{y}^N = A^N \mathbf{y}^0 + \sum_{i=0}^{N-1} A^i B \mathbf{u}^{i(N-1)}$$



**Initial state parameterisation
does not work with inputs!**

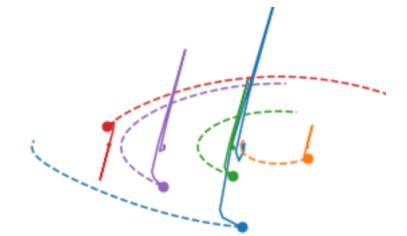
Question: Did we learn a method to parameterise the Koopman operator for brownian motion?



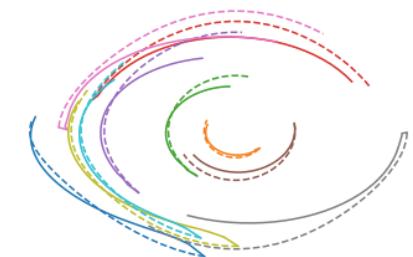
DENIS+, w Brownian inputs



Trained DENIS+ on 0 inputs



DENIS+, w Brownian + 0 inputs

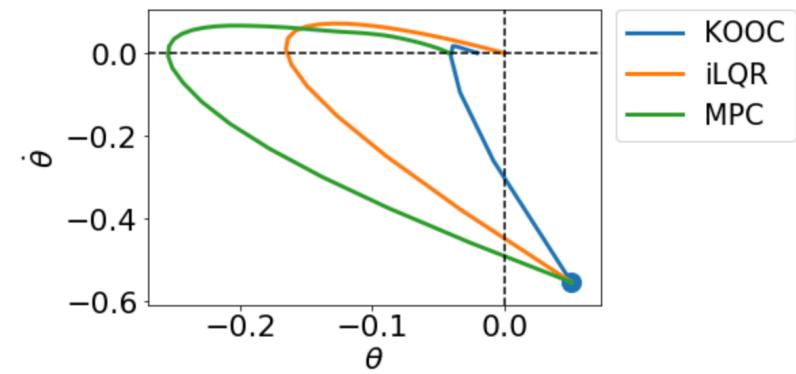
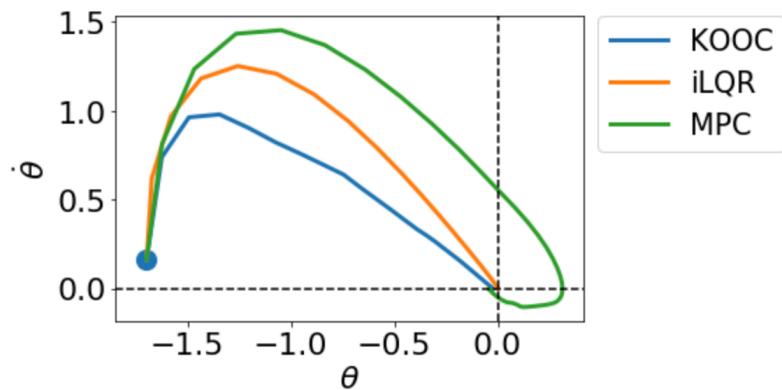
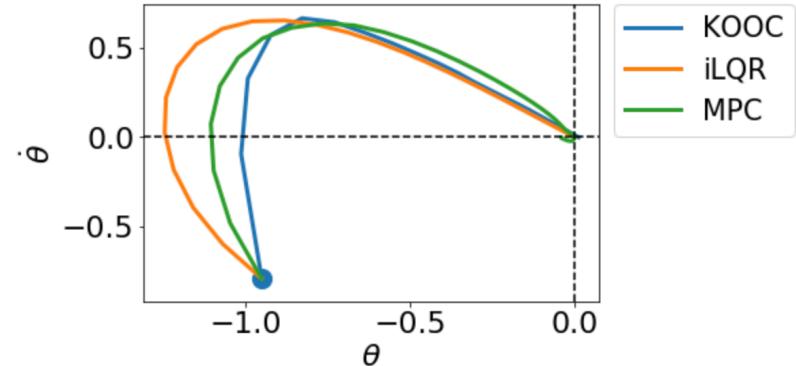
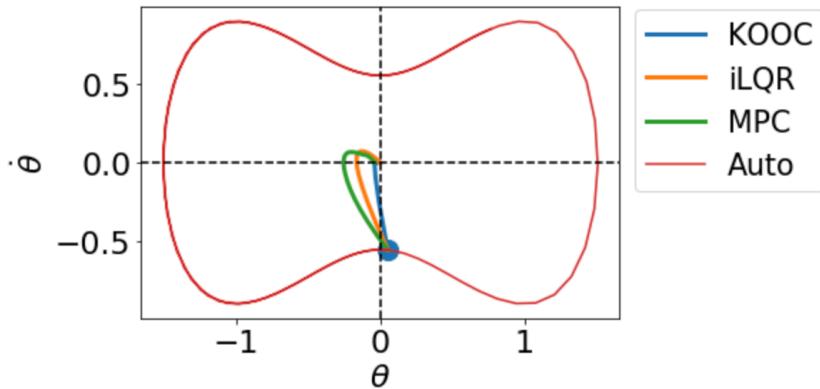


Project Summary W8-9

DENIS

Once Upon a Time...

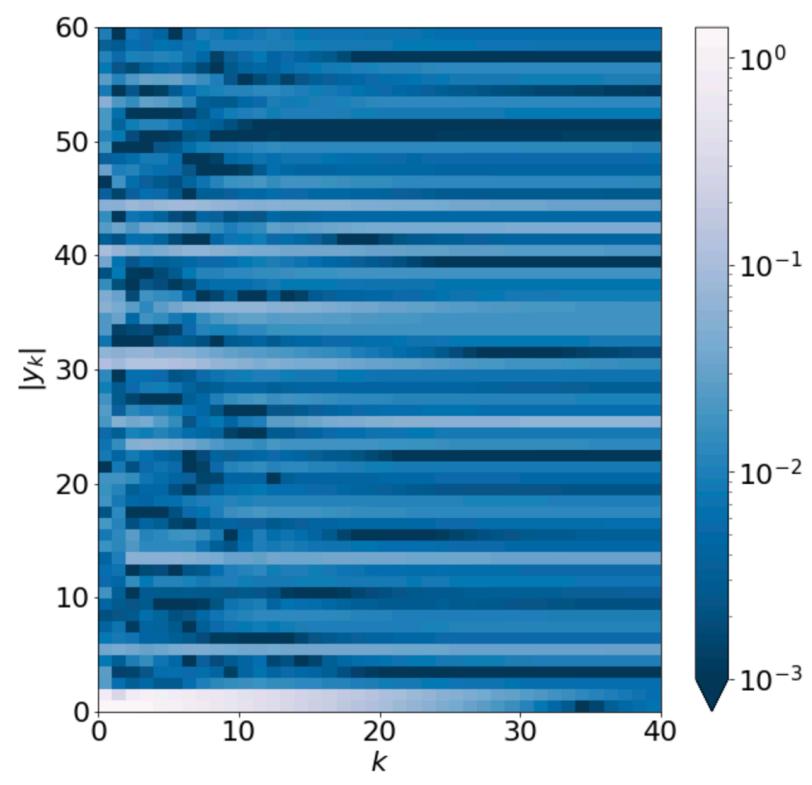
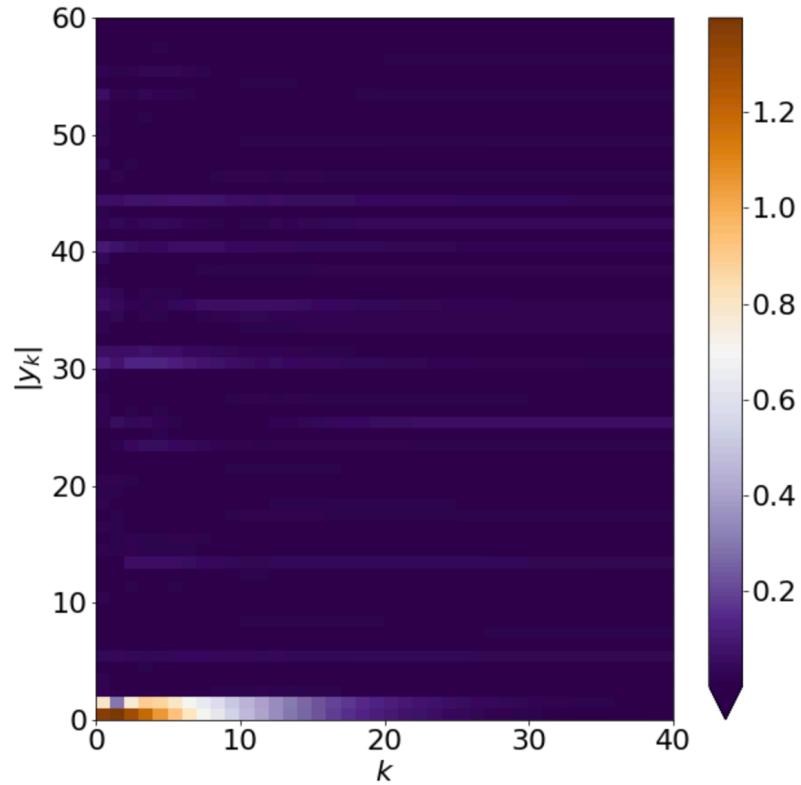
- Bug discovered with DENIS KOOC.... Now KOOC beats iLQR (I have been using the same koopman operator for every initial conditions!)



DENIS

Once Upon a Time...

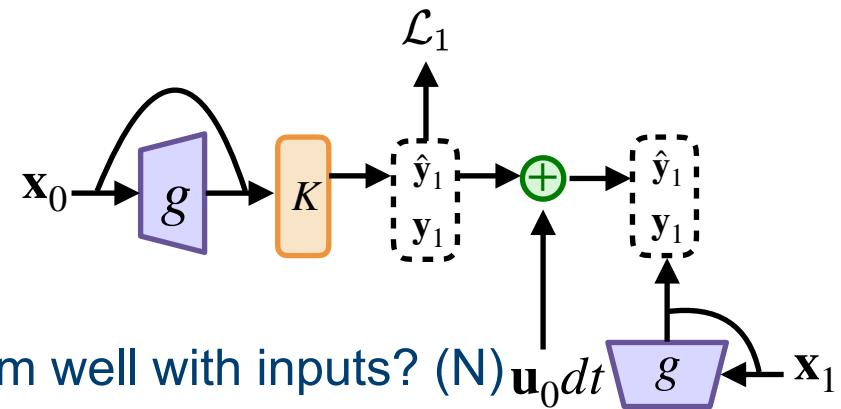
- How about the latent space?



LREN

Sanity check for input affine systems

- We have shown, at large dt , LREN cannot be modified to include inputs due to instability. At small dt , does it perform well? (Y)



- Does LREN trained without inputs perform well with inputs? (N) $\mathbf{u}_0 dt$

$$\mathbf{y}^N = \mathbf{A}^N \mathbf{y}^0 + \sum_{i=0}^{N-1} \mathbf{A}^i \mathbf{B} \mathbf{u}^{i(N-1)}$$

- Does LREN trained with inputs perform well without inputs? (Y)

LREN

Data Generation

- Some time dedicated to write a general ‘data generator’ method (a necessary evil)

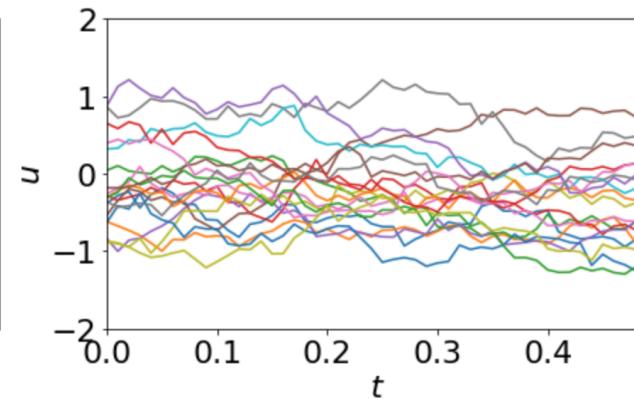
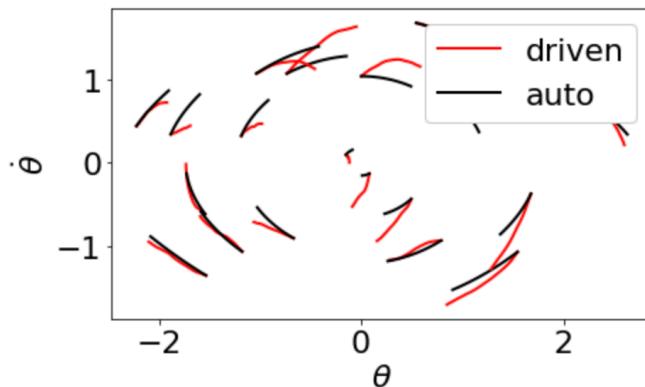
```
N = 25000
n_steps = 50
dt = 0.01

generator = Data_Generator(pendulum_dynamics, dt, 2)
```

- Inputs are now drawn from OU processes

$$X_{n+1} = X_n + \theta (\mu - X_n) \Delta t + \sigma \Delta W_n$$

$$\mu \sim U(-1, 1), \quad \sigma = 1, \quad \theta = 2$$



LREN

Results

A: Autonomous trajectories, no input

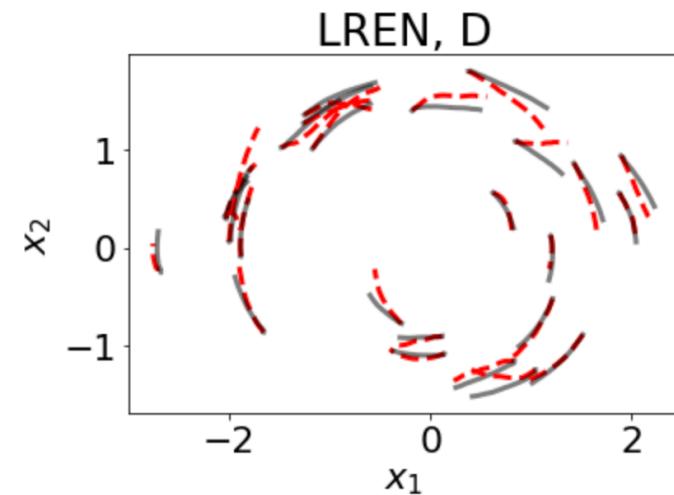
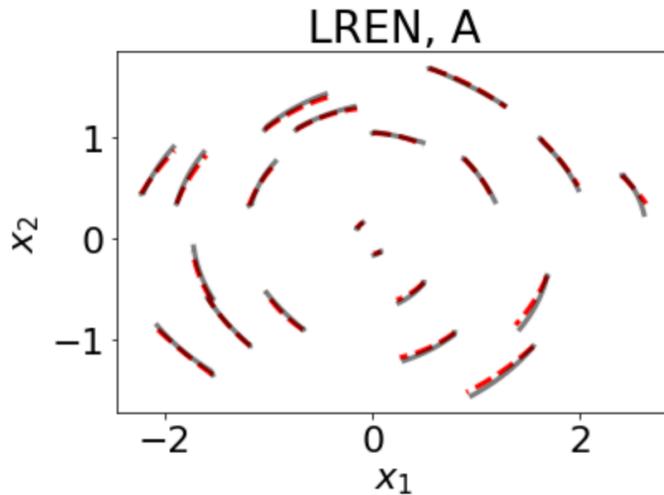
D: Driven trajectories, with input

LREN: Base LREN model

LREN + : LREN model, with added input during training, trained with balanced A and D trajectories

LREN - : Same as previous, but only trained on D trajectories

- Does LREN trained without inputs perform well with inputs?



LREN

Results

A: Autonomous trajectories, no input

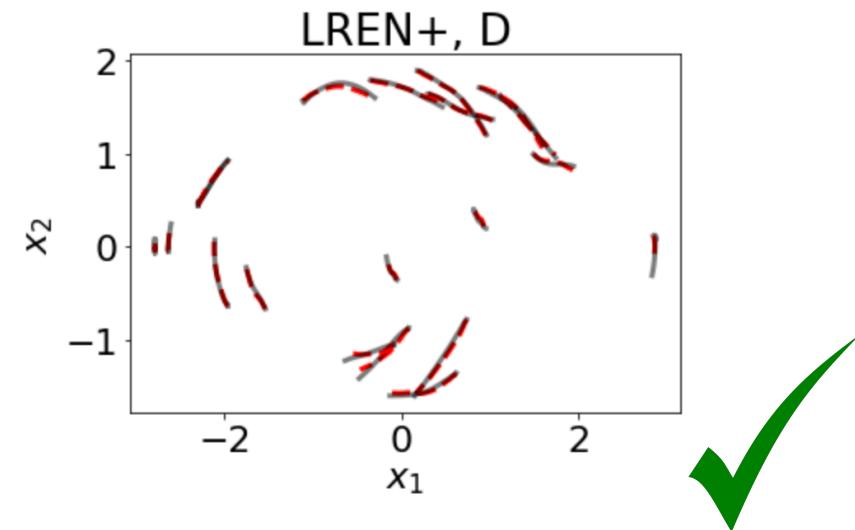
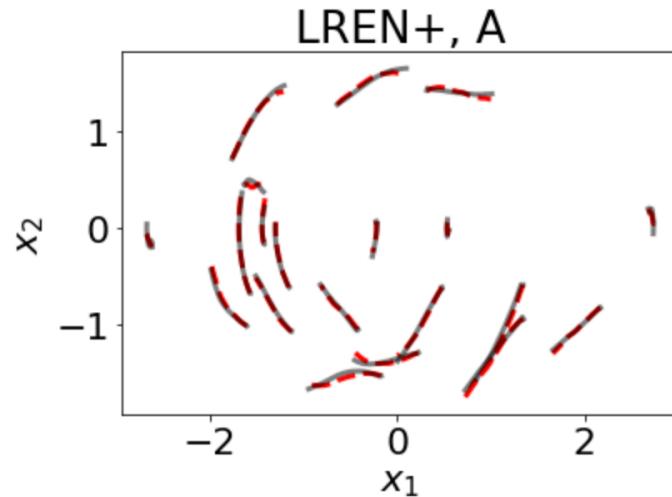
D: Driven trajectories, with input

LREN: Base LREN model

LREN + : LREN model, with added input during training, trained with balanced A and D trajectories

LREN - : Same as previous, but only trained on D trajectories

- What if we allow LREN model to train with inputs? (A+D Dataset)



LREN

Results

A: Autonomous trajectories, no input

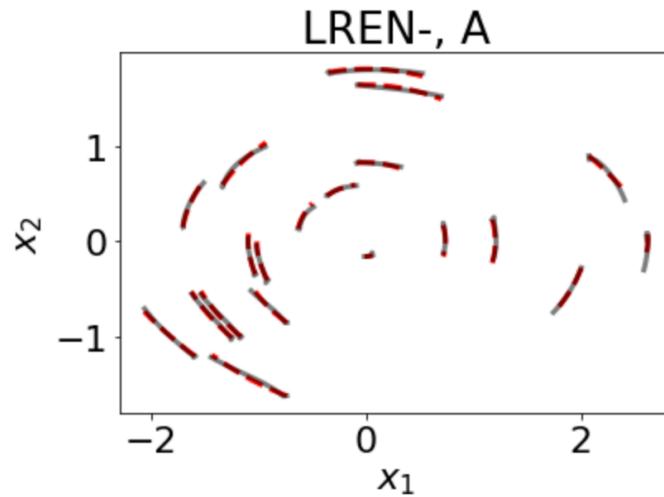
D: Driven trajectories, with input

LREN: Base LREN model

LREN + : LREN model, with added input during training, trained with balanced A and D trajectories

LREN - : Same as previous, but only trained on D trajectories

- What if the model is trained only on driven data? Can it reconstruct autonomous trajectories?



LREN

Results

A: Autonomous trajectories, no input

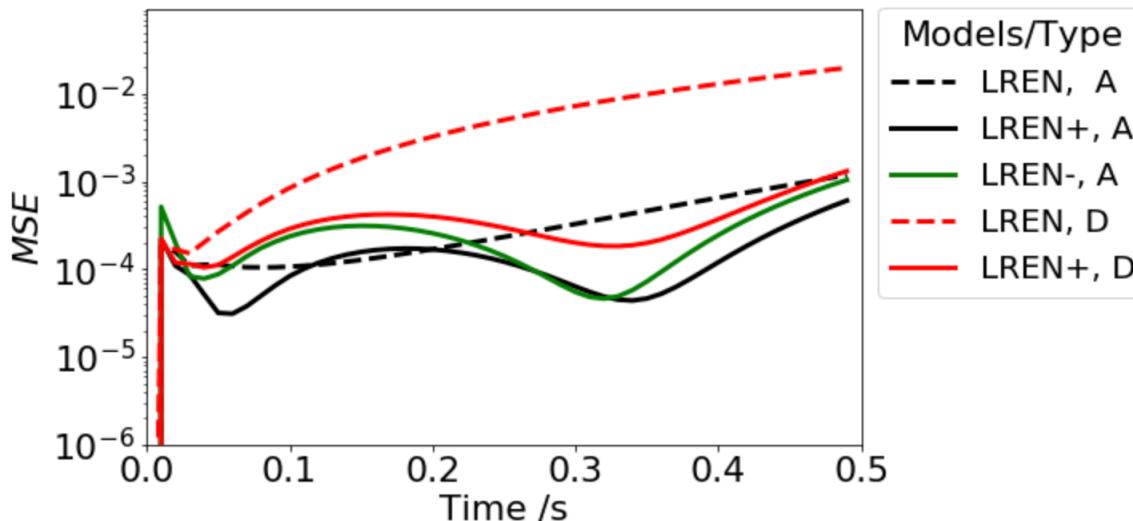
D: Driven trajectories, with input

LREN: Base LREN model

LREN + : LREN model, with added input during training, trained with balanced A and D trajectories

LREN - : Same as previous, but only trained on D trajectories

- What about overall performance across all methods investigated?

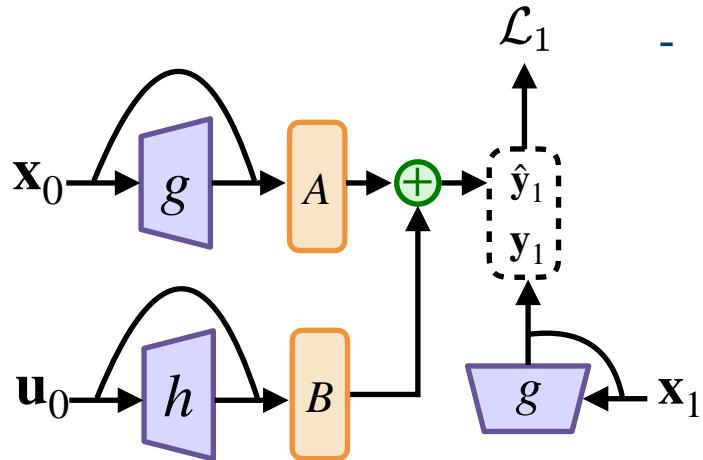


- At SMALL time steps:
- Training without inputs will not predict driven trajectories
- Training with inputs may even improve autonomous performance when dataset is balanced
- Training with only driven trajectories may slightly hurt autonomous prediction

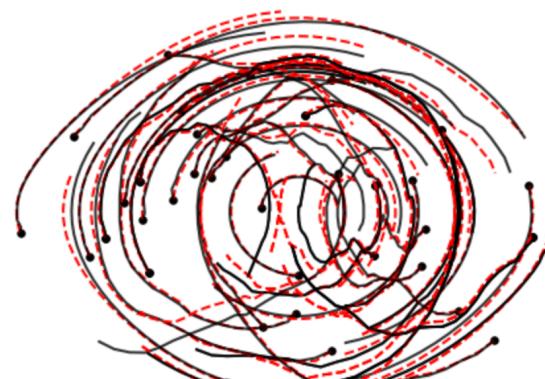
DEINA

Double Encoder for Input Non-Affine Systems

- Bug fix, architecture is confirmed

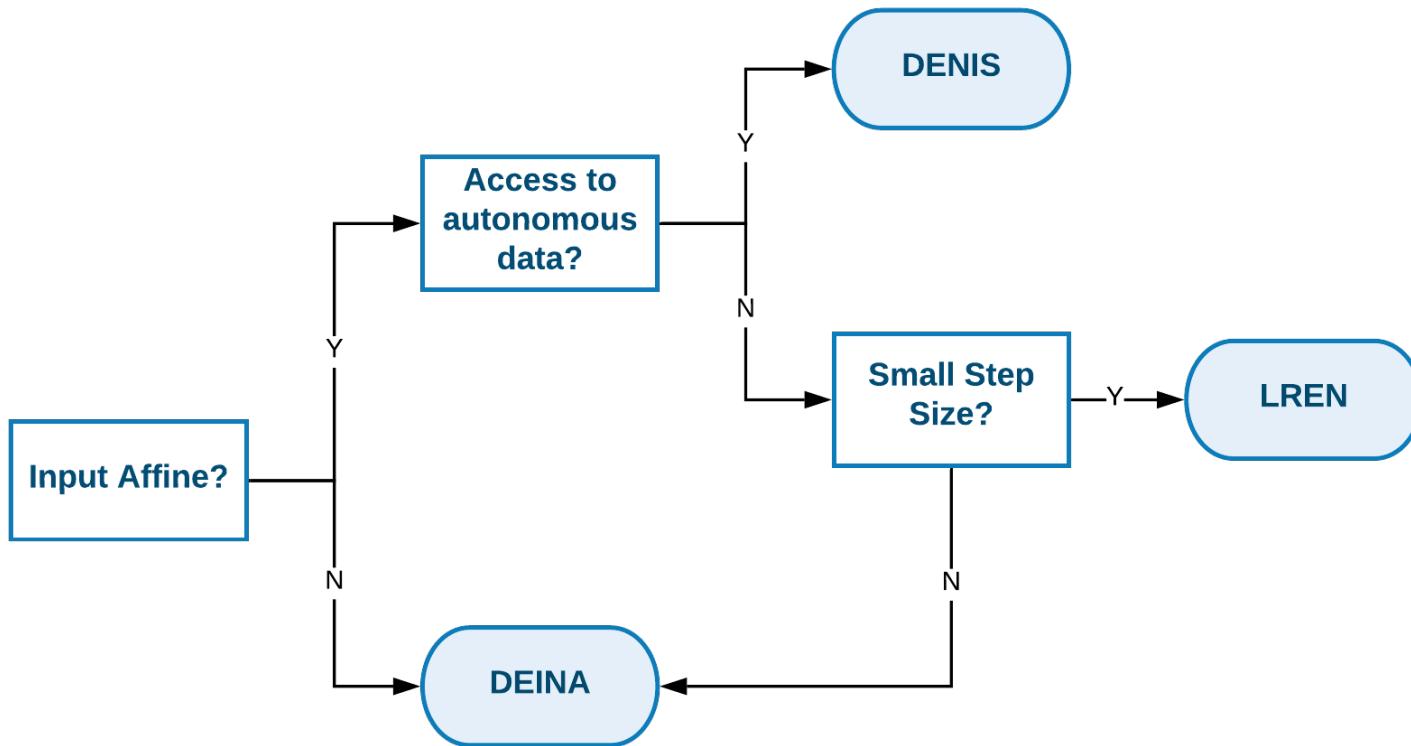


- Why do we need to append input?
 - Without it, we will need to include a decoder (MPC, LQR will not work as we do not know how to reconstruct the latent inputs)
 - Because of the ‘many to one’ issue (Rory), the decoder must be linear
 - A linear decoder is equivalent to appending input



:)

DL Koopman Decision Tree:



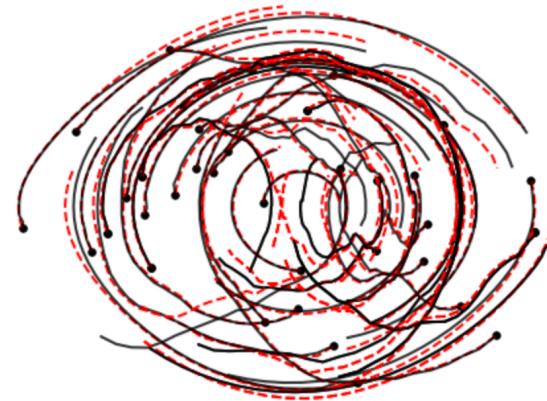
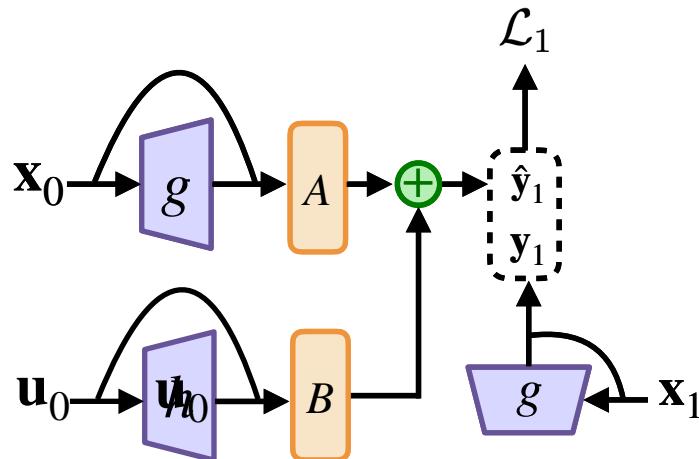
Project Summary

Week ∞

DEINA DEINA DEINA

DEINA: First Iteration

- $\mathbf{x} \in \mathbb{R}^N, \mathbf{u} \in \mathbb{R}^M, A \in \mathbb{R}^{\hat{N} \times \hat{N}}, \dots$
- B operates on $[\mathbf{u}, h(\mathbf{u})]^T$ but is optimal control obvious/sensible?
- i.e. For the LQR gain matrix $G \in \mathbb{R}^{\hat{N} \times \hat{N}}$, can we simply ditch the columns to get $G \in \mathbb{R}^{\hat{N} \times M}$?
- Or.... Shall we just ditch h all together?

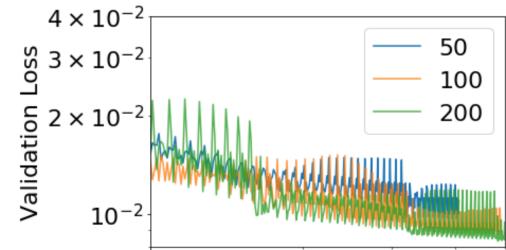
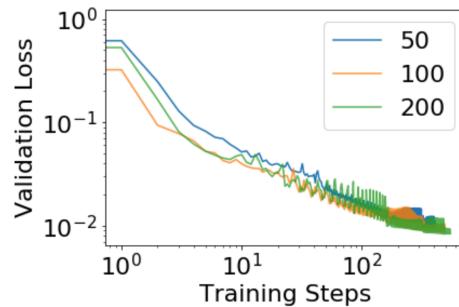


RMSE: 0.094 c.f. 0.024, 30TH@0.1s

DEINA: First Iteration

- So.... Results not great
- Increase model capacity? Fix encoder dimension to be $2 \times \frac{N}{4} \times \frac{N}{2} \times \frac{N}{2} \times N$
- LR Scheduler: $lr^*=0.5$ per 20 epochs

| N | Epochs | RMSE |
|-----|--------|-------|
| 50 | 50 | 0.090 |
| 100 | 40 | 0.089 |
| 200 | 80 | 0.085 |
| 300 | 100 | 0.085 |



Work's focus: make DEINA better

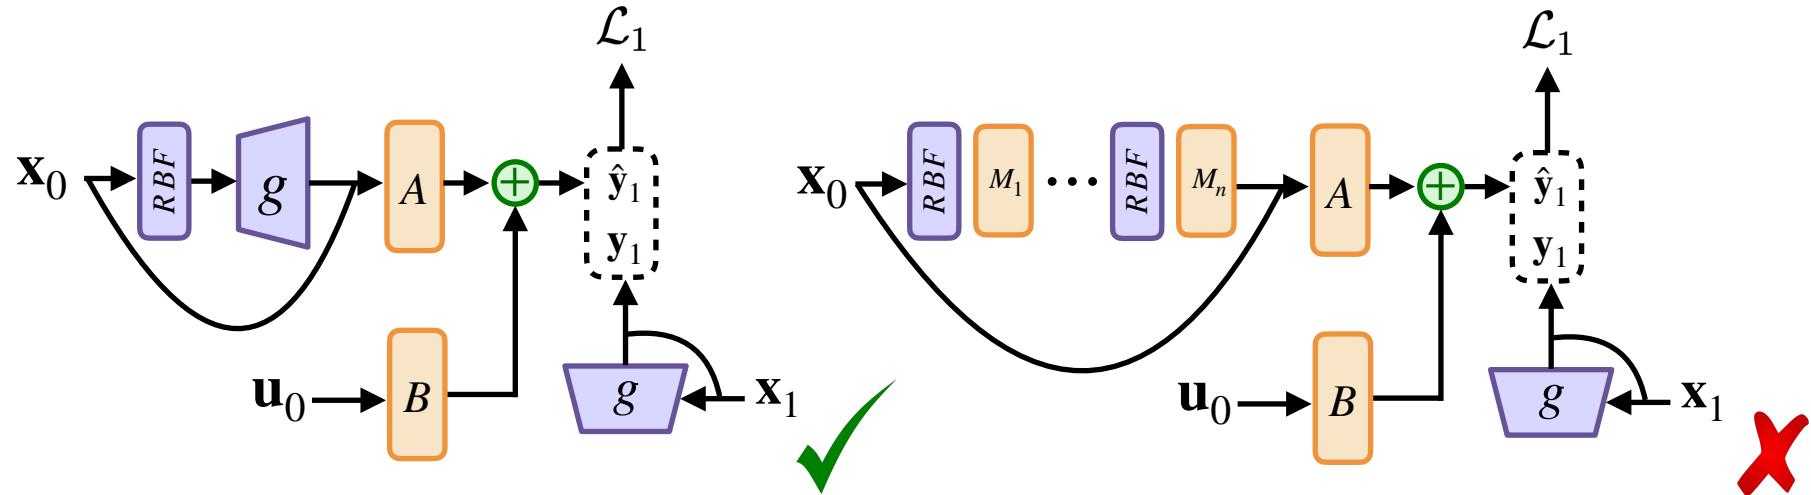
- It appears size doesn't matter?? How?

DEINA RBF

- Is there another way to increase the capacity of the network other than upping the latent?
- Make the model more “non-parametric”
- RBF Networks?

$$\begin{bmatrix} \hat{\mathbf{x}}_1 \\ \vdots \\ \hat{\mathbf{x}}_n \end{bmatrix} = \begin{bmatrix} \phi\left(\sigma_1 \|\mathbf{x}_1 - \mathbf{c}_1\|_2\right) & \cdots & \phi\left(\sigma_m \|\mathbf{x}_1 - \mathbf{c}_m\|_2\right) \\ \vdots \\ \phi\left(\sigma_1 \|\mathbf{x}_n - \mathbf{c}_1\|_2\right) & \cdots & \phi\left(\sigma_m \|\mathbf{x}_n - \mathbf{c}_m\|_2\right) \end{bmatrix}$$

| N | Epochs | RMSE |
|------------------|--------|--------------|
| 50 | 50 | 0.090 |
| 100 | 40 | 0.089 |
| 200 | 80 | 0.085 |
| 300 | 100 | 0.085 |
| RBF - 100 | 50 | 0.071 |



DEINA ODE

- Something fun to do....
- Motivation: Maybe, we will get better result if the Koopman operator is continuous. Maybe, we can recreate DENIS behaviour (bounded long TH trajectories). No need for discrete LQR

Discrete Dynamics

$$\mathbf{y} = g(\mathbf{x}_{k+1})$$

$$\begin{bmatrix} \mathbf{x}_{k+1} \\ \mathbf{y}_{k+1} \end{bmatrix} = K \begin{bmatrix} \mathbf{x}_k \\ \mathbf{y}_{k+1} \end{bmatrix}$$

Continuous Dynamics

$$\dot{\mathbf{y}}_{t_1} = K\mathbf{y}_{t_1}$$

$$\mathbf{y}_{t_2} = \mathbf{y}_{t_1} + \int_{t_1}^{t_2} \dot{\mathbf{y}}_{t_1} dt$$

$$\mathbf{x}_{t_2} = C\mathbf{y}_{y_2}$$

Neural Ordinary Differential Equations

Ricky T. Q. Chen*, Yulia Rubanova*, Jesse Bettencourt*, David Duvenaud
 University of Toronto, Vector Institute
 {rtqichen, rubanova, jessebett, duvenaud}@cs.toronto.edu

Backprop = Solving another ODE, but backwards in time!

$$\mathbf{a}(t) = \partial L / \partial \mathbf{y}(t) \quad \mathbf{s}_0 = \left[\mathbf{y}(t_1), \frac{\partial L}{\partial \mathbf{y}(t_1)}, \mathbf{0} \right] \quad \nabla = \begin{bmatrix} Ky(t) \\ -\mathbf{a}(t)^\top K \\ -\mathbf{a}(t)^\top \frac{\partial K}{\partial \theta} \end{bmatrix}$$

$$\left[\mathbf{y}(t_0), \frac{\partial L}{\partial \mathbf{y}(t_0)}, \frac{\partial L}{\partial \theta} \right] = \text{ODESolve}(\mathbf{s}_0, \nabla, t_1, t_0, \theta)$$

DEINA ODE

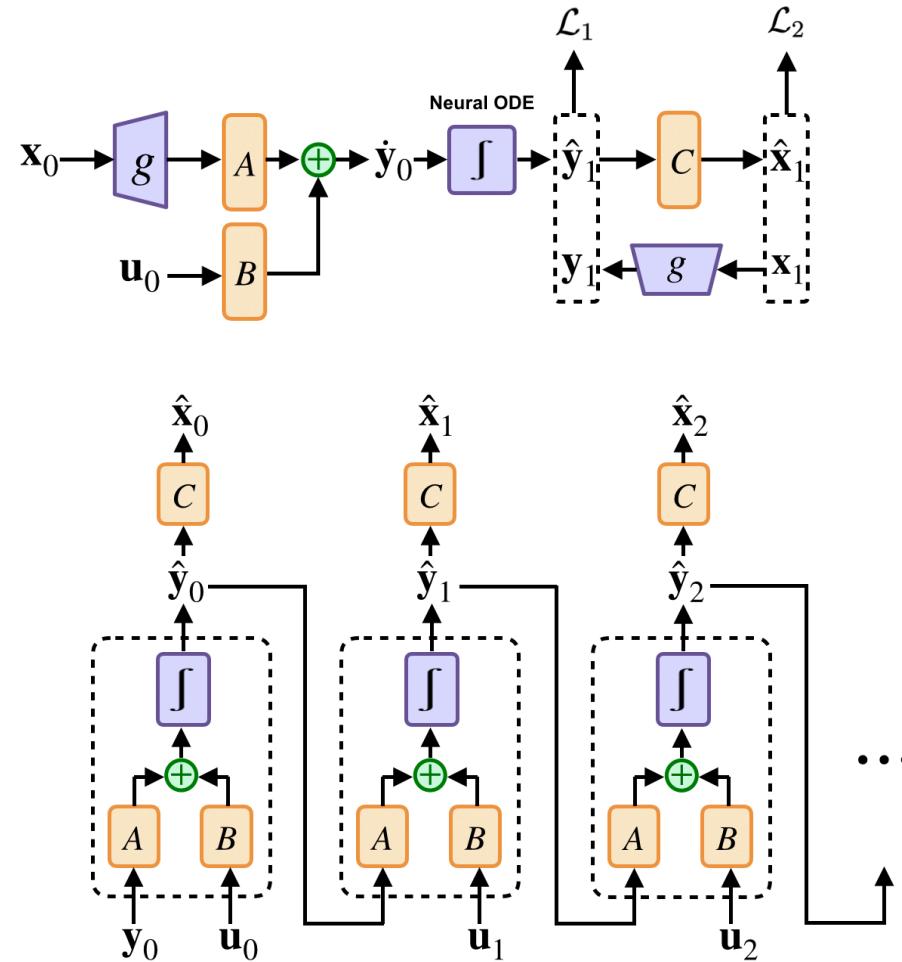
$$\dot{\mathbf{y}}_{t_1} = A\mathbf{y}_{t_1} + B\mathbf{u}_{t_1}$$

$$\mathbf{y}_{t_2} = \mathbf{y}_{t_1} + \int_{t_1}^{t_2} \dot{\mathbf{y}}_{t_1} dt$$

$$\mathbf{x}_{t_2} = C\mathbf{y}_{t_2}$$

Initial Evaluation:

- (+) Converges within 2-5 epochs!
- (+) Smooth trajectories
- (-) RMSE does not improve, if not slightly worse than before :(



Summary & Future Work

- DEINA, DEINA RBF, DEINA ODE presented
- DEINA RBF does increase model performance
- DEINA ODE is more like a proof of concept i,e. Continuous dynamics can be learnt using discrete data.
 - Are there more benefits to it?
 - Does it outperform previous methods with KOOC? Need to evaluate
 - Can MSE be improved with a RBF layer?



Summary & Future Work

- **Tidy up code, environment, write actual training script (for HP search), etc**
- Analyse these methods in more depth- MSE VS TH plots, etc
- Is KOOC with  possible? Does it actually bring any benefit with model performance?
- Arm model, might be time?

